

Container

A standard unit of software
that packages up code and
all its dependencies
so the application runs
quickly and reliably
from one computing
environment to another.

- Docker



Container

A standard unit of software
that packages up code and
all its dependencies
so the application runs
quickly and reliably
from one computing
environment to another.

```
• FROM bowl:1.0.0
• RUN apt-get install \
    h2o \
    pebbles \
    plant
• COPY fish .
• CMD ["fish", "swim"]
```

- Docker



Container

A standard unit of software
that packages up code and
all its dependencies
so the application runs
quickly and reliably
from one computing
environment to another.

```
• FROM bowl:1.0.0
• RUN apt-get install \
    h2o \
    pebbles \
    plant
• COPY fish .
• CMD ["fish", "swim"]
```

- Docker



- Design time: Cohesive full stack of apps and dependencies
- Runtime: Boxed process on an operating system

Container

A standard unit of software
that packages up code and
all its dependencies
so the application runs
quickly and reliably
from one computing
environment to another.

```
• FROM bowl:1.0.0
  • RUN apt-get install \
    h2o \
    pebbles \
    plant
  • COPY fish .
  • CMD ["fish", "swim"]
```

- Docker

- Design time: Cohesive full stack of apps and dependencies
- Runtime: Boxed process on an operating system



Why Containers?

Why would you make your life more complicated?

- Build once, run anywhere
- Polyglot
- Disposable
- Modular
- Separation of concerns
- Own your infrastructure
- Rapid scaling

Why Containers?

Why would you make your life more complicated?

Build once, run anywhere

Polyglot

Disposable

Modular

Separation of concerns

Own your infrastructure

Rapid scaling

Replication:

Production abhors singletons

Why Containers?

Why would you make your life more complicated?

Build once, run anywhere

Polyglot

Disposable

Modular

Separation of concerns

Own your infrastructure

Rapid scaling

Replication:

Production abhors singletons

New reality for service delivery

Why Containers?

Why would you make your life more complicated?

- Build once, run anywhere
- Polyglot
- Disposable
- Modular
- Separation of concerns
- Own your infrastructure
- Rapid scaling

Replication:
Production abhors singletons

New reality for service delivery



THE TWELVE-FACTOR APP

12factor.net

Why Containers?

Why would you make your life more complicated?

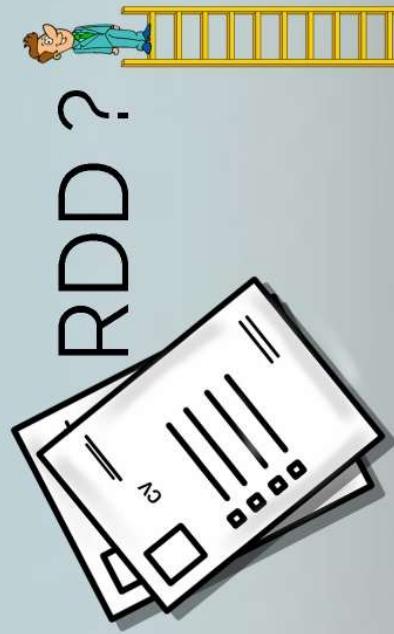
- Build once, run anywhere
- Polyglot
- Disposable
- Modular
- Separation of concerns
- Owning your infrastructure
- Rapid scaling

Replication:
Production abhors singletons

New reality for service delivery



12factor.net



Not related to Spark
Resilient Distributed Datasets (RDDs)

Layers



Infrastructure as code

Layers



Code
Functions
Logic

Infrastructure as code

Layers

Dependencies

Code
Functions
Logic



Infrastructure as code

Layers

Package

Dependencies

Code
Functions
Logic



Infrastructure as code

Layers

API / Web Service

- REST

- gRPC

- GraphQL

Package

Dependencies

Code
Functions
Logic



Infrastructure as code

Layers

API / Web Service

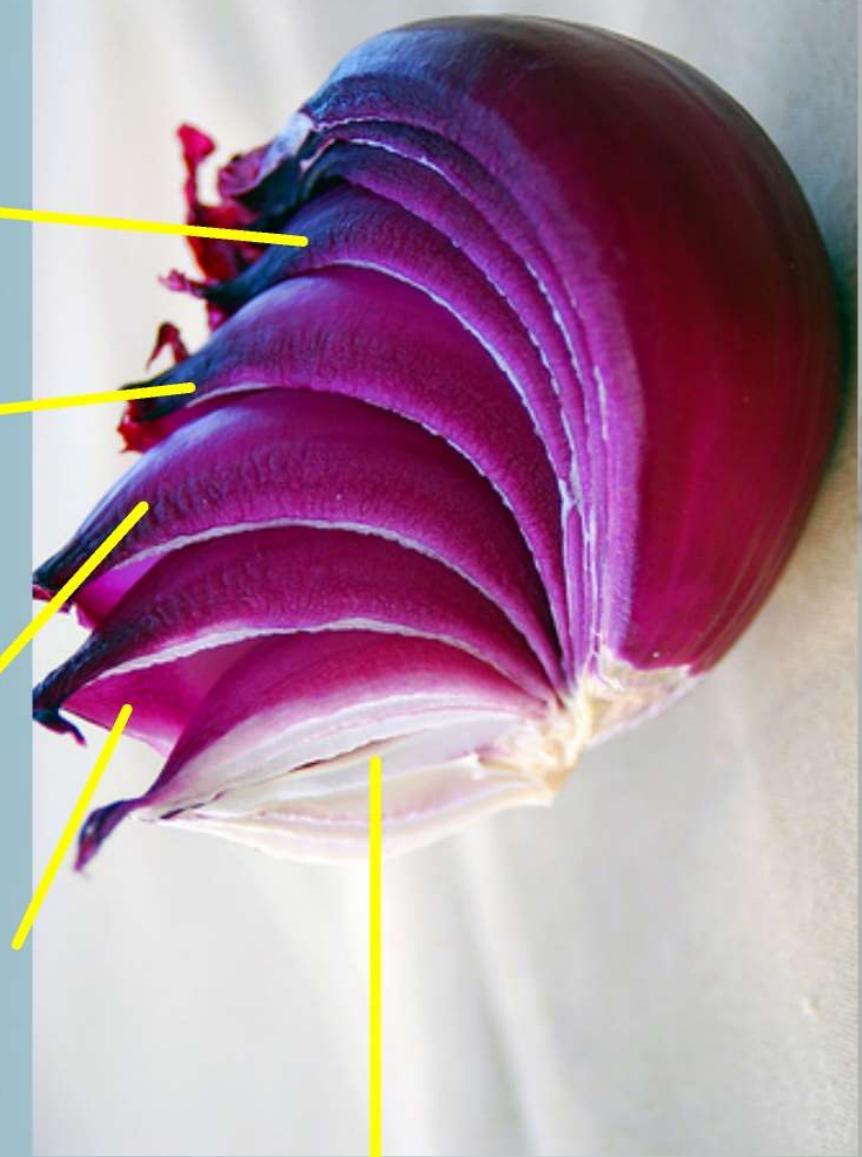
- REST

- gRPC

- GraphQL

Package

Dependencies

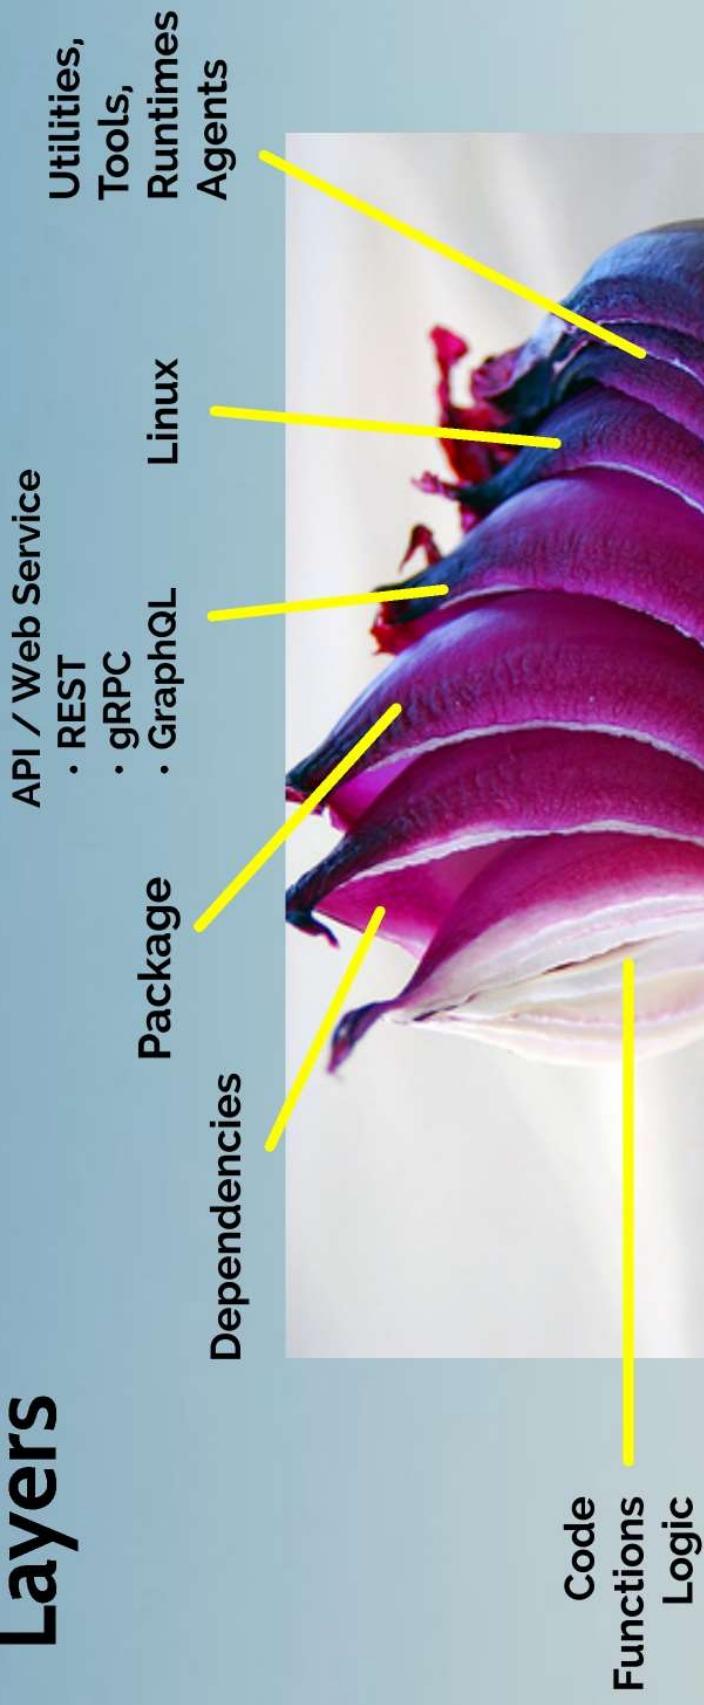


Code
Functions
Logic

Linux

Infrastructure as code

Layers



Infrastructure as code

Layers

API / Web Service

- REST
- gRPC
- GraphQL

Package

Dependencies

Linux

Utilities,
Tools,
Runtimes
Agents

Code
Functions
Logic

Container
Image



Infrastructure as code

Layers

API / Web Service

- REST
- gRPC
- GraphQL

Package

Dependencies

Linux

Utilities,
Tools,
Runtimes
Agents

Code
Functions
Logic

Container
Image

Name:Version

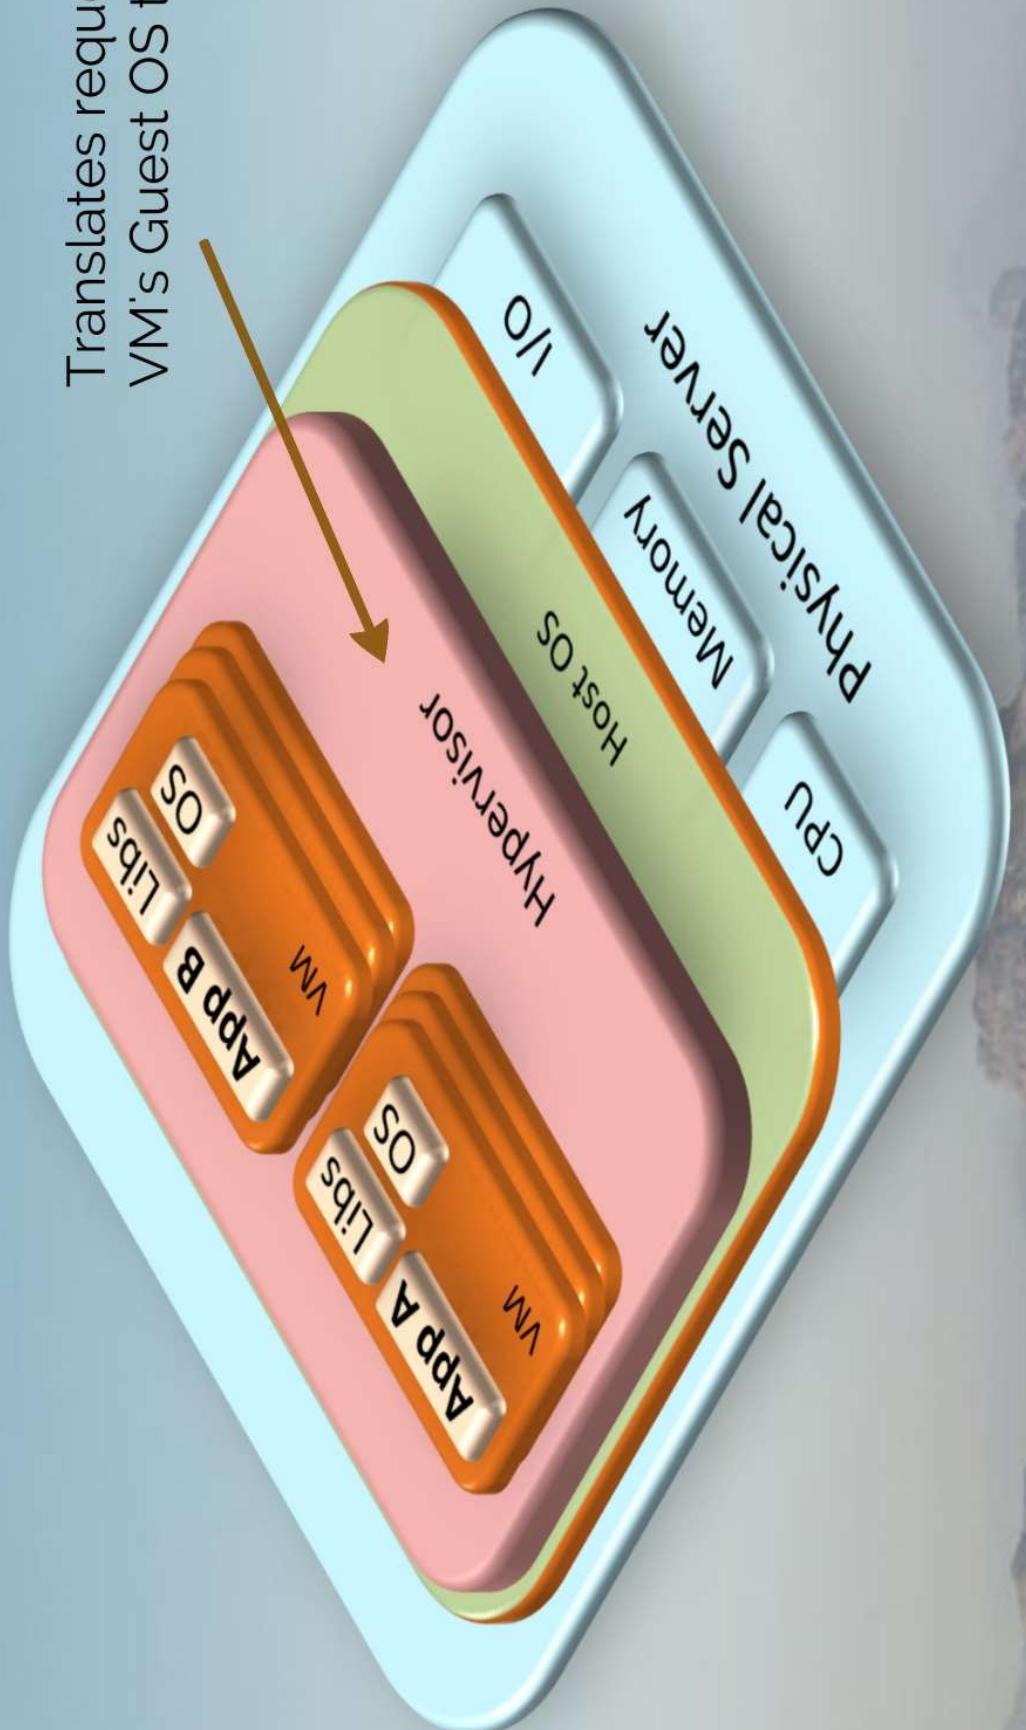


Registries



Infrastructure as code

Before Containers, We had Virtual Machines



Container Architecture

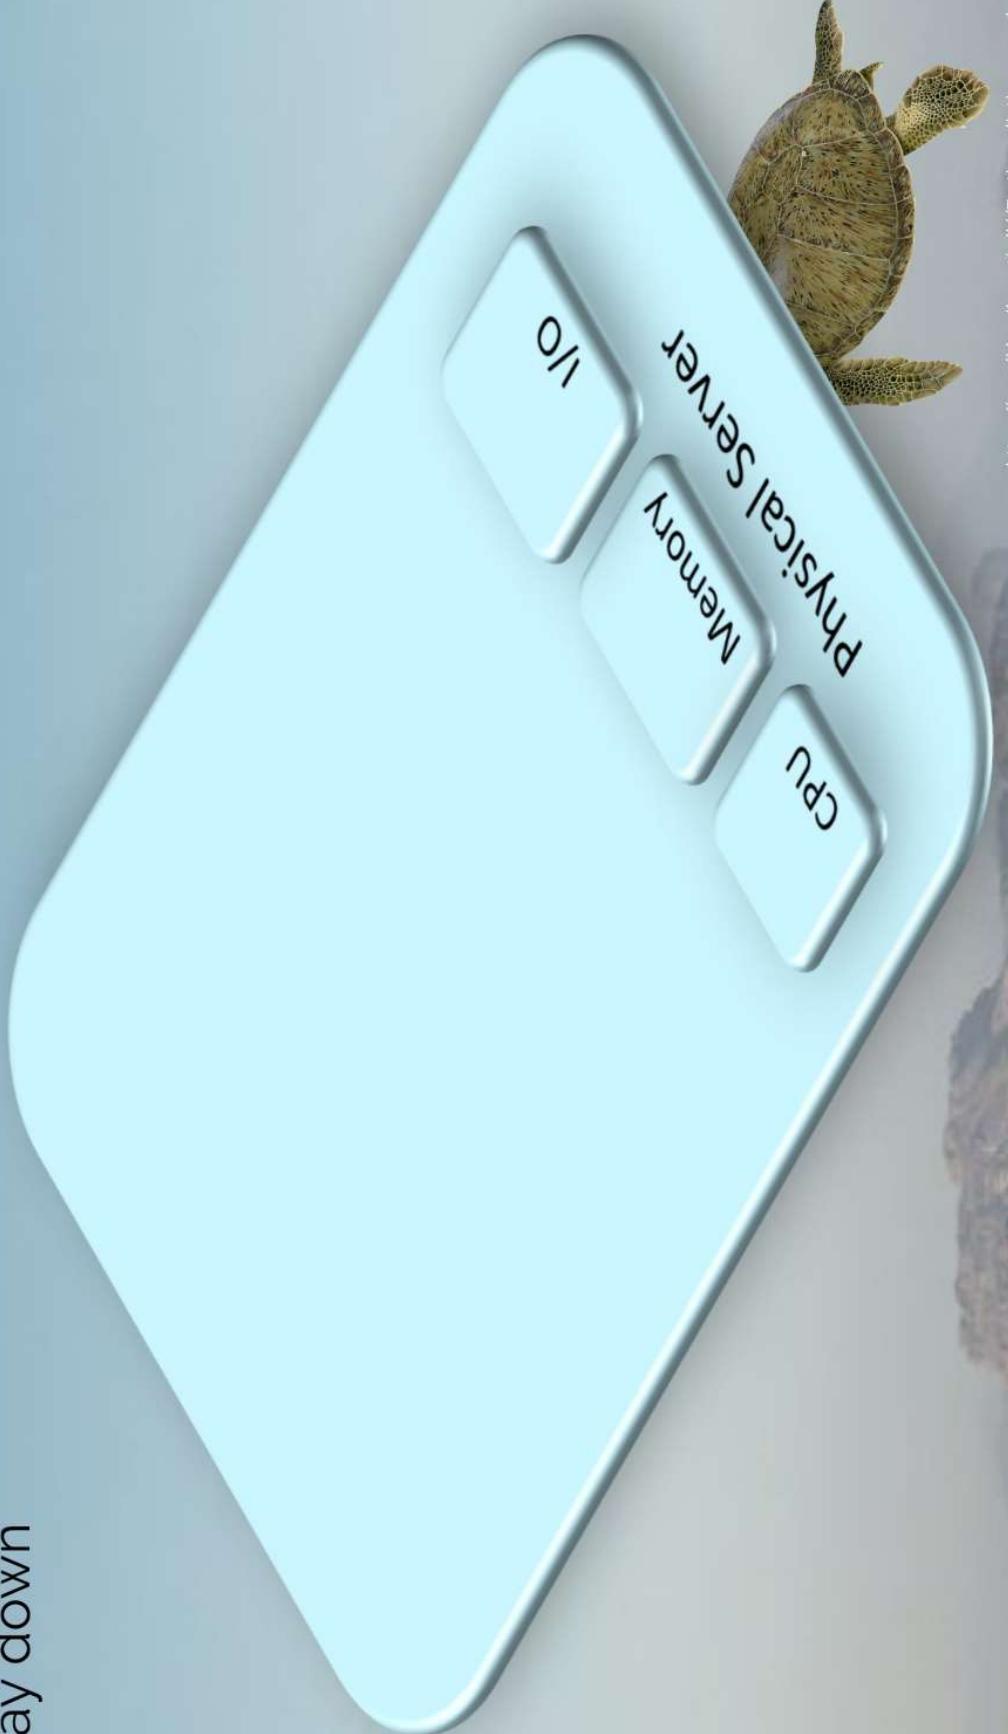
Turtles all the way down



https://en.wikipedia.org/wiki/Turtles_all_the_way_down

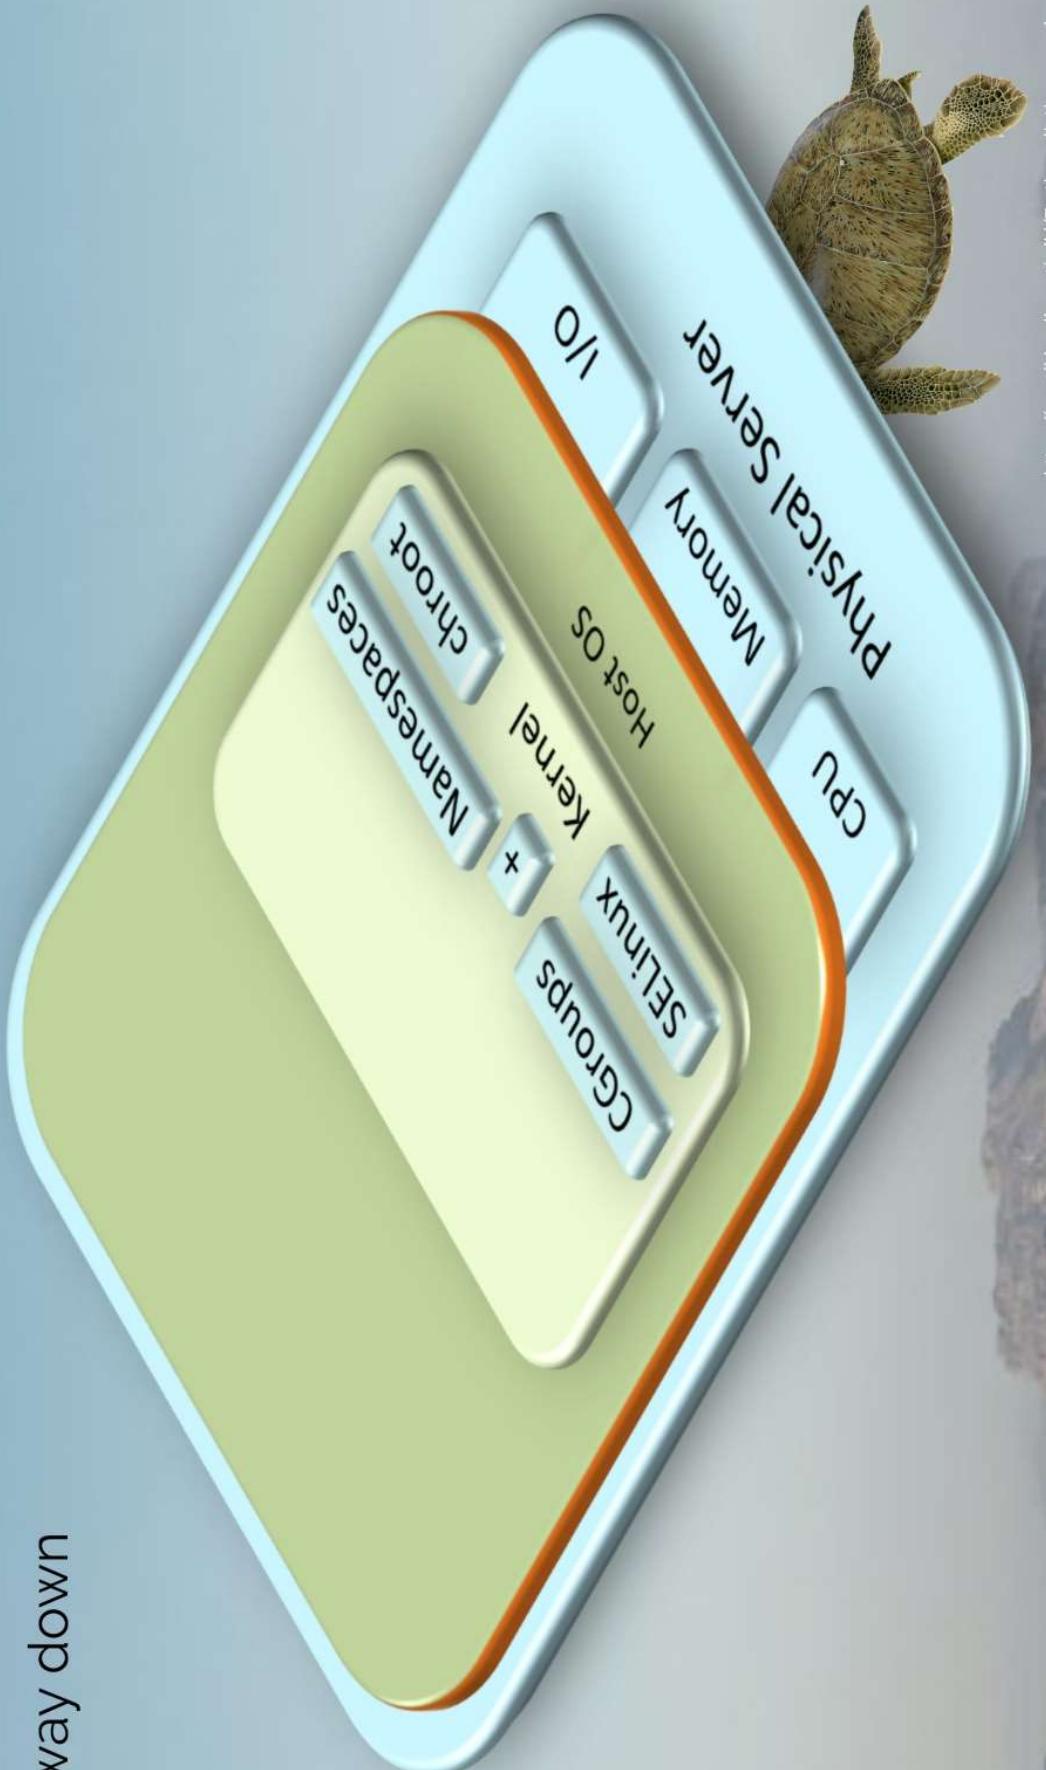
Container Architecture

Turtles all the way down



Container Architecture

Turtles all the way down



Container Architecture

Turtles all the way down

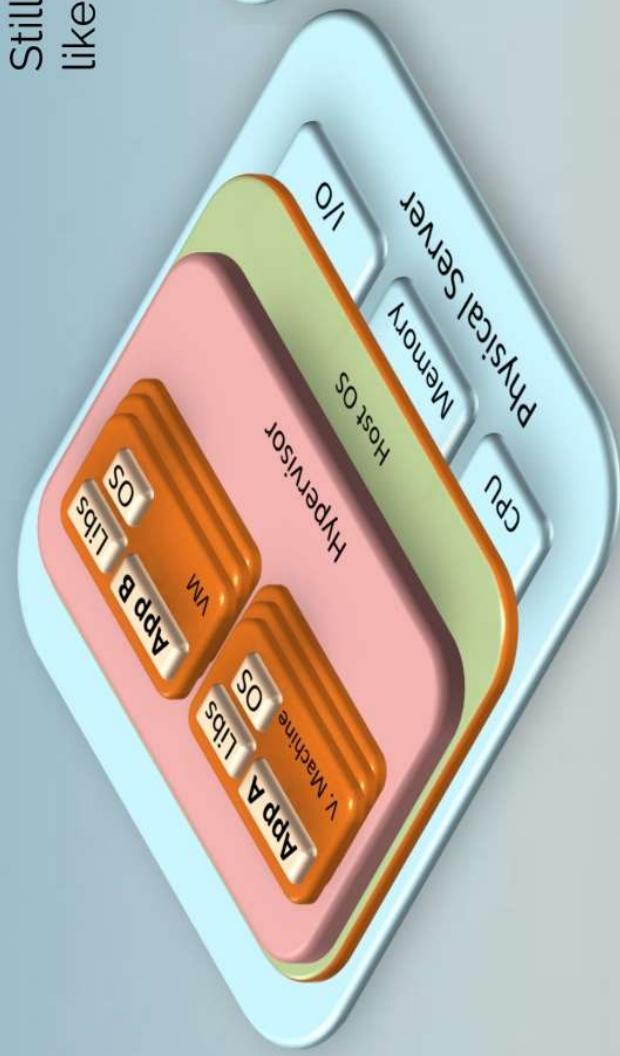
OS in container
delegates to host
OS kernel services



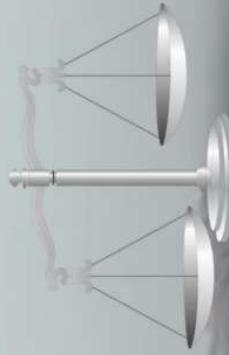
Host OS securely
divvies-up resources to
each isolated container

Evolution: VMs

to Container Runtimes



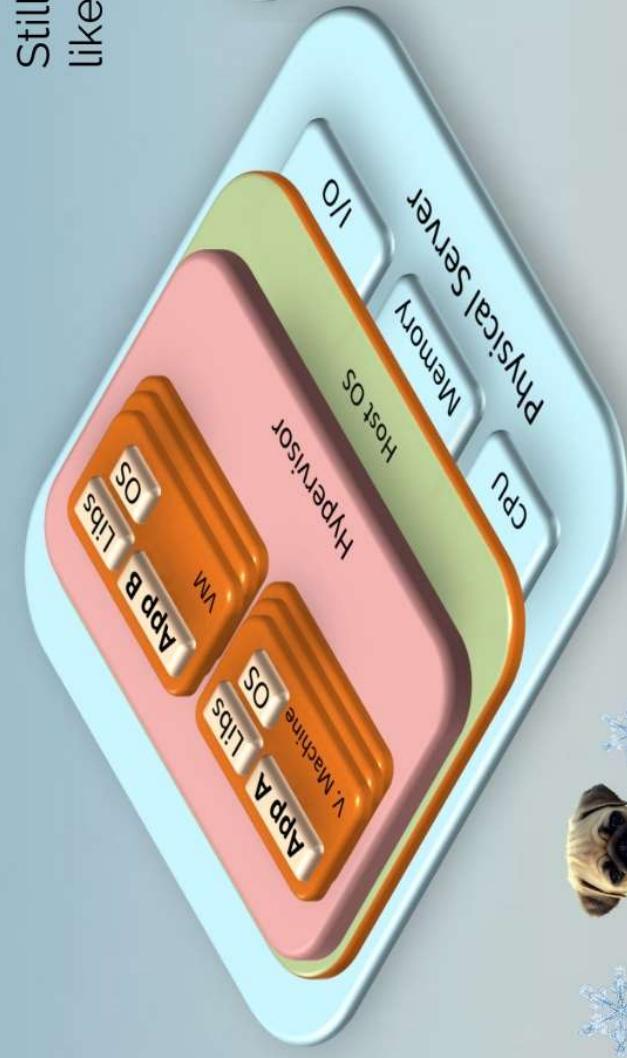
Still feels
like a VM



Evolution: VMs

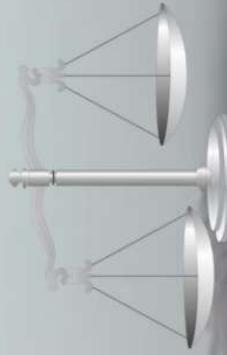
to Container Runtimes

Still feels
like a VM



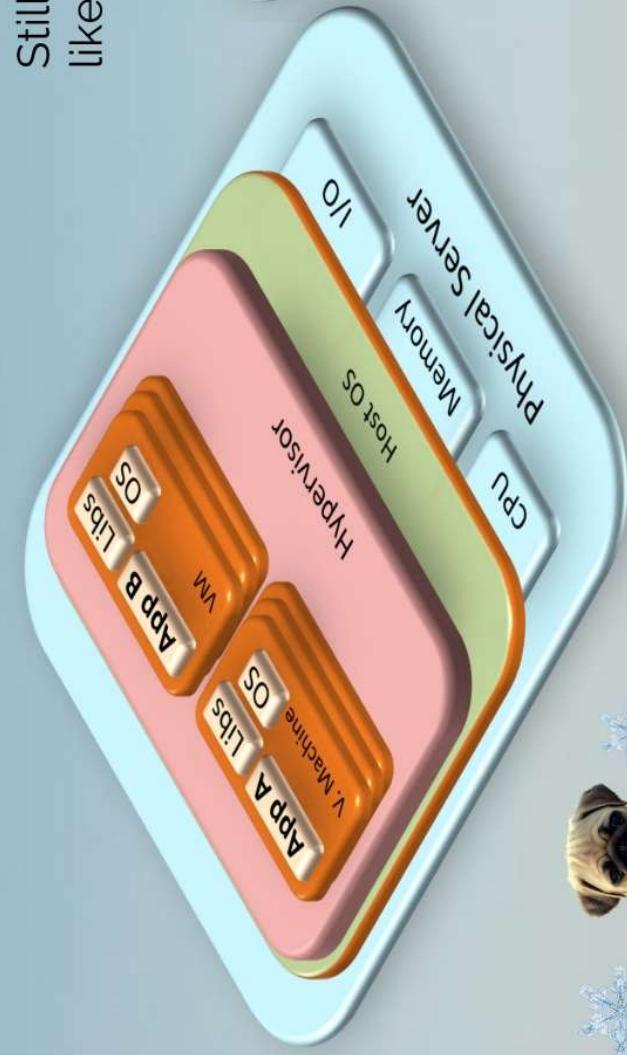
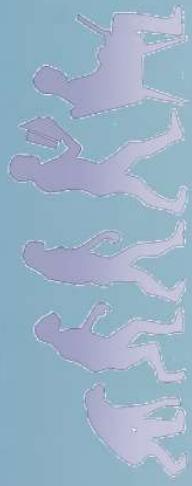
Unique servers,
snowflakes GBs -
minutes to start

<https://martinfowler.com/bliki/SnowflakeServer.html>

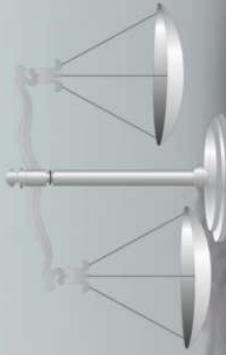


Evolution: VMs

to Container Runtimes



Disposable, ephemeral IPs
MBs - seconds to start



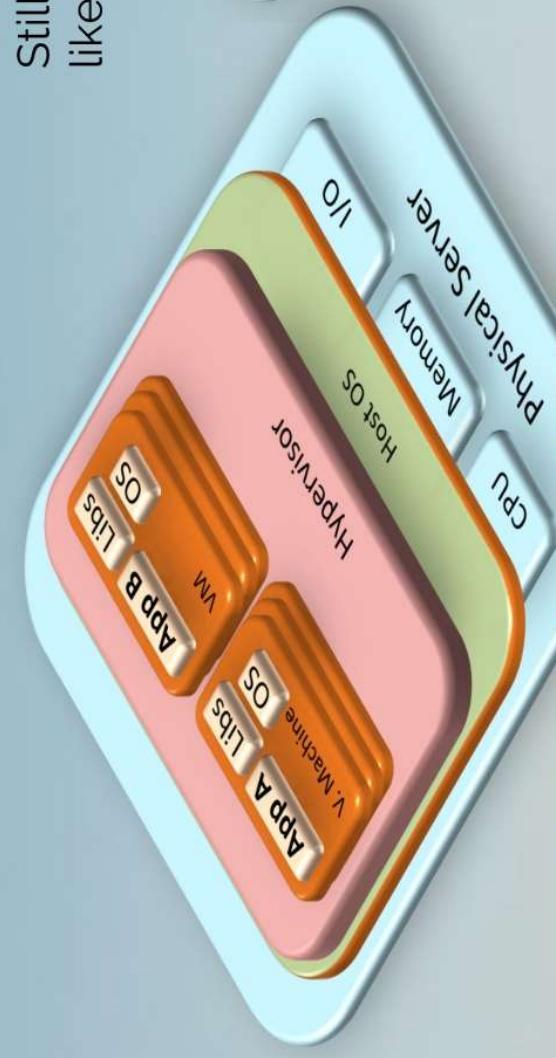
Unique servers,
snowflakes GBs -
minutes to start



Evolution: VMs

to Container Runtimes

Still feels
like a VM



Kernel Sleight of Hand for Containers



Kernel Sleight of Hand for Containers

chroot

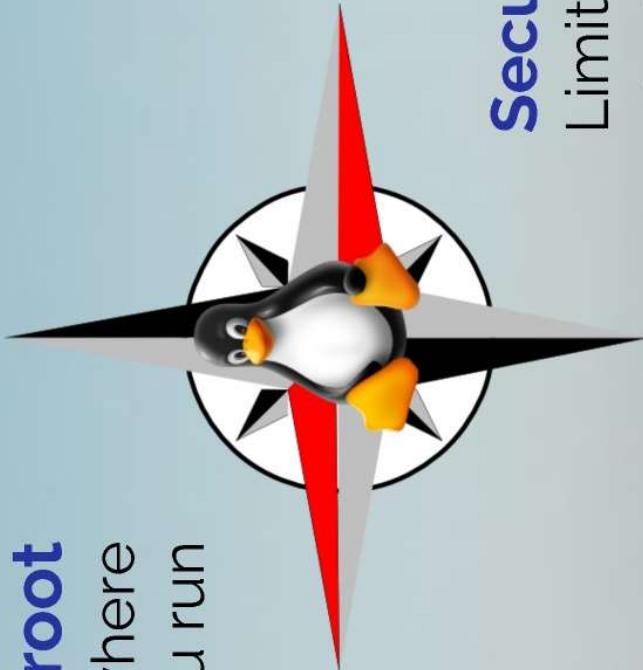
Limit where
you run



Kernel Sleight of Hand for Containers

chroot

Limit where
you run



Security-Enhanced Linux

Limit what
you can access

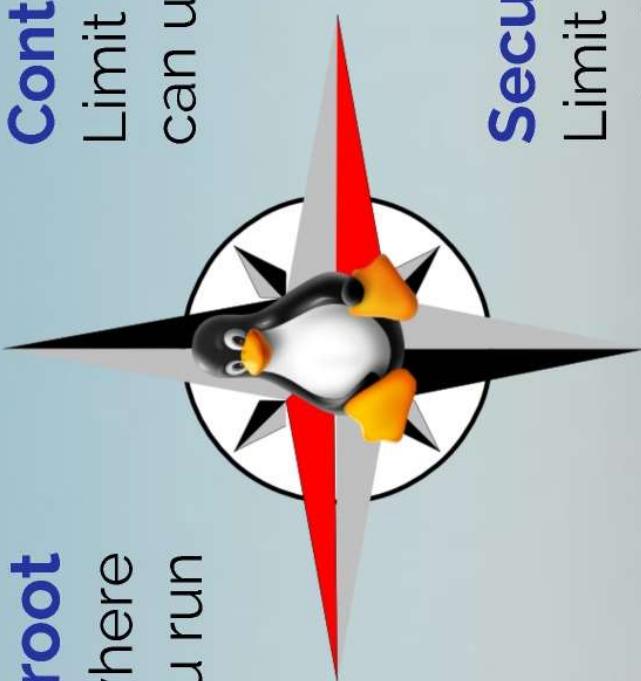
Kernel Sleight of Hand for Containers

chroot

Limit where
you run

Control Groups

Limit what you
can use



Security-Enhanced Linux

Limit what
you can access

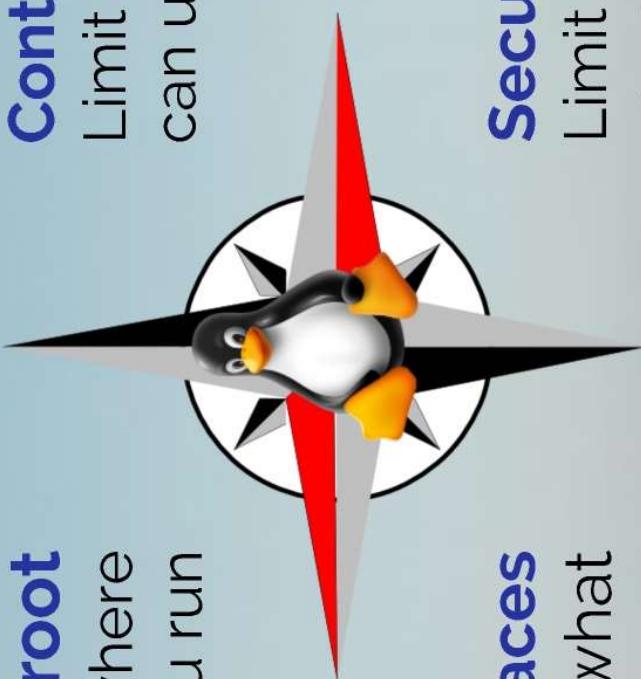
Kernel Sleight of Hand for Containers

chroot

Limit where
you run

Control Groups

Limit what you
can use



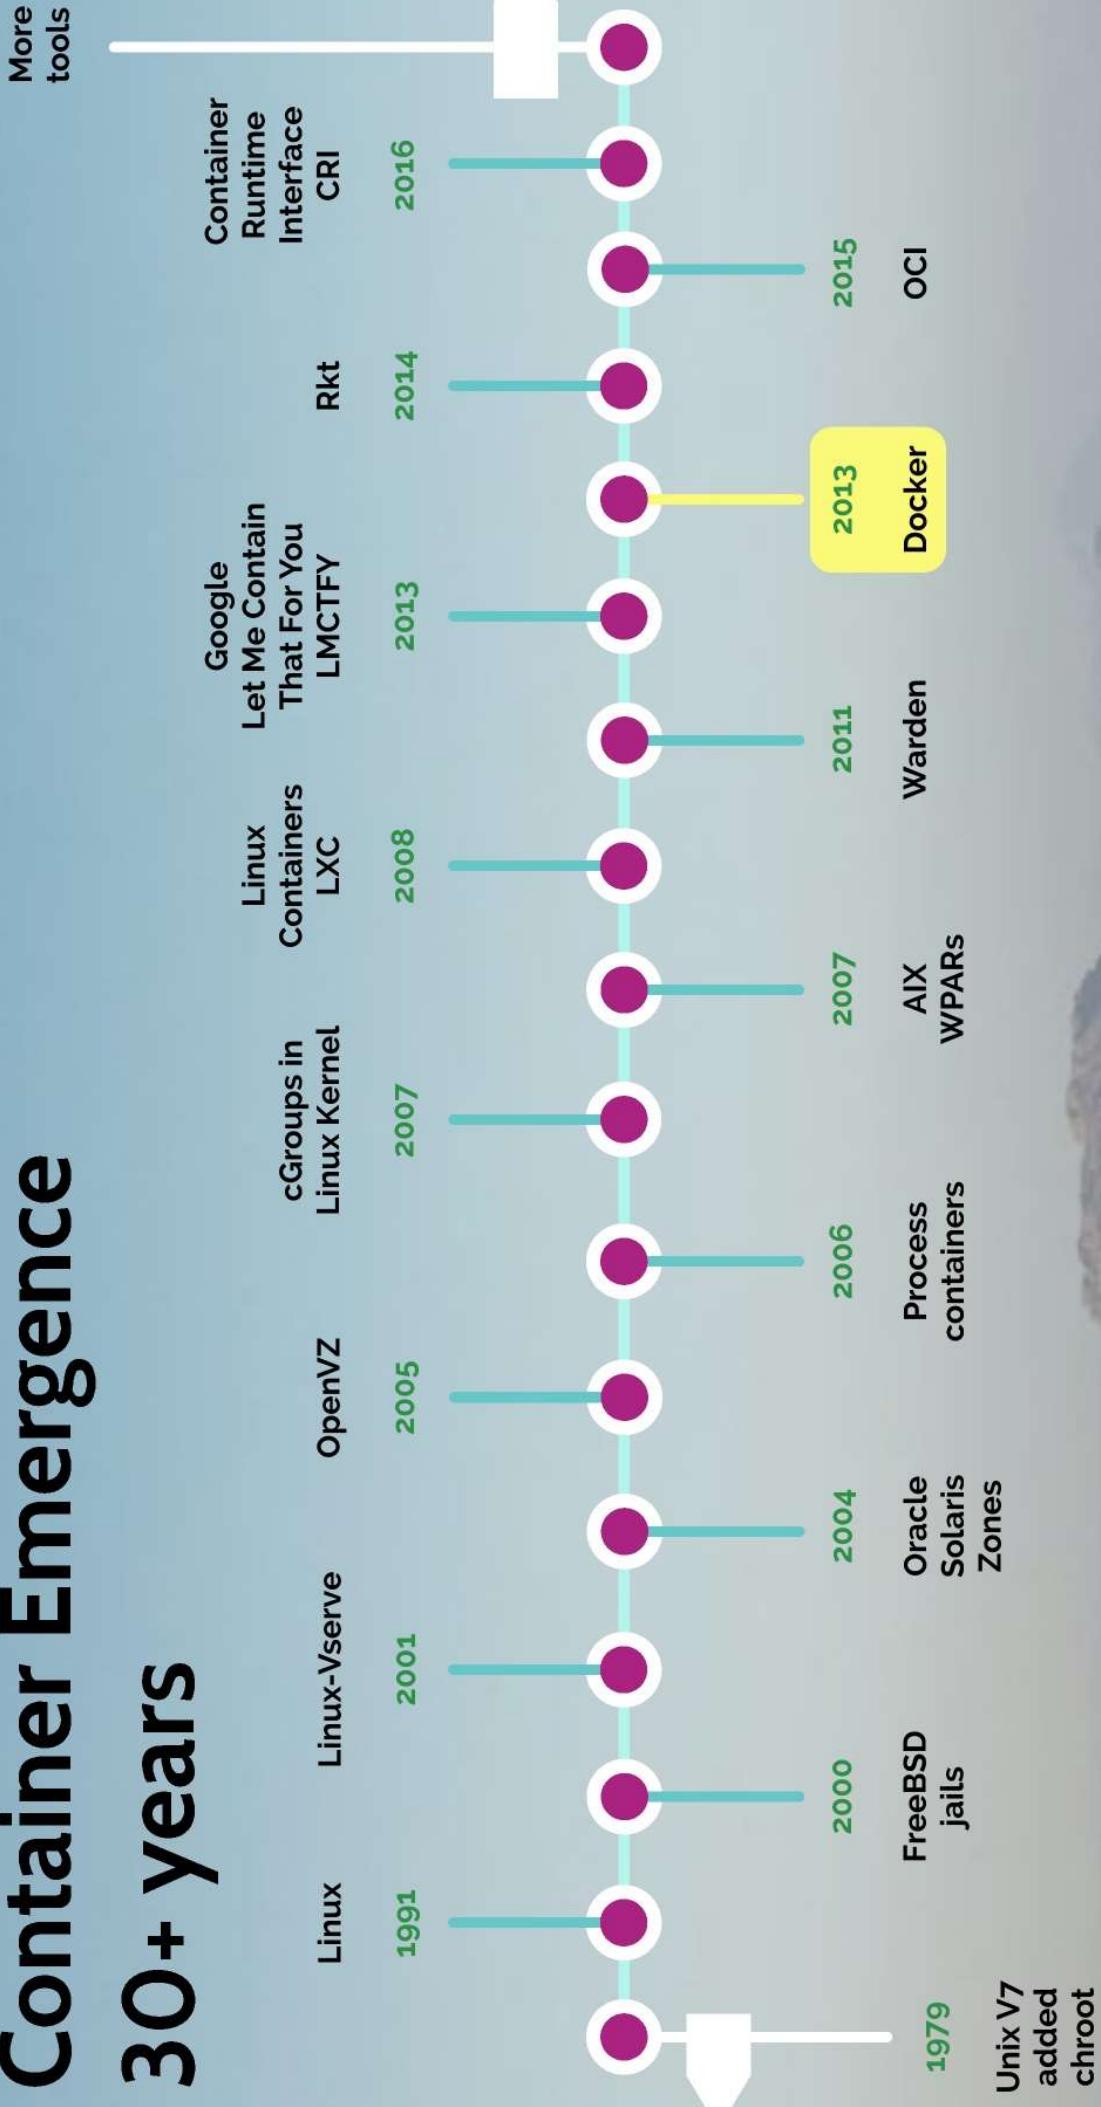
Namespaces

Limit what
you can see

Security-Enhanced Linux

Limit what
you can access

Container Emergence 30+ years



Terminology

A few common terms

Container Image

A file of bytes following the container image format that defines the container. Stored in a repository. Often composed of multiple layers.

Image Layer

Container images are typically many layers of images. Each layer defines distinct and reusable container definitions. FROM, COPY, ADD and RUN add new layers

Container Engine

Process on the Host OS that accepts requests to pull images, calls the container runtime and cleans up after containers terminate.

Registry

Library or datastore of multiple repositories. Hosts and manages the repo files and manifests (add, remove, list, security).

Repository

Listing of one or more versions of a container along with its images layers and manifests. Follows established image manifest schema.

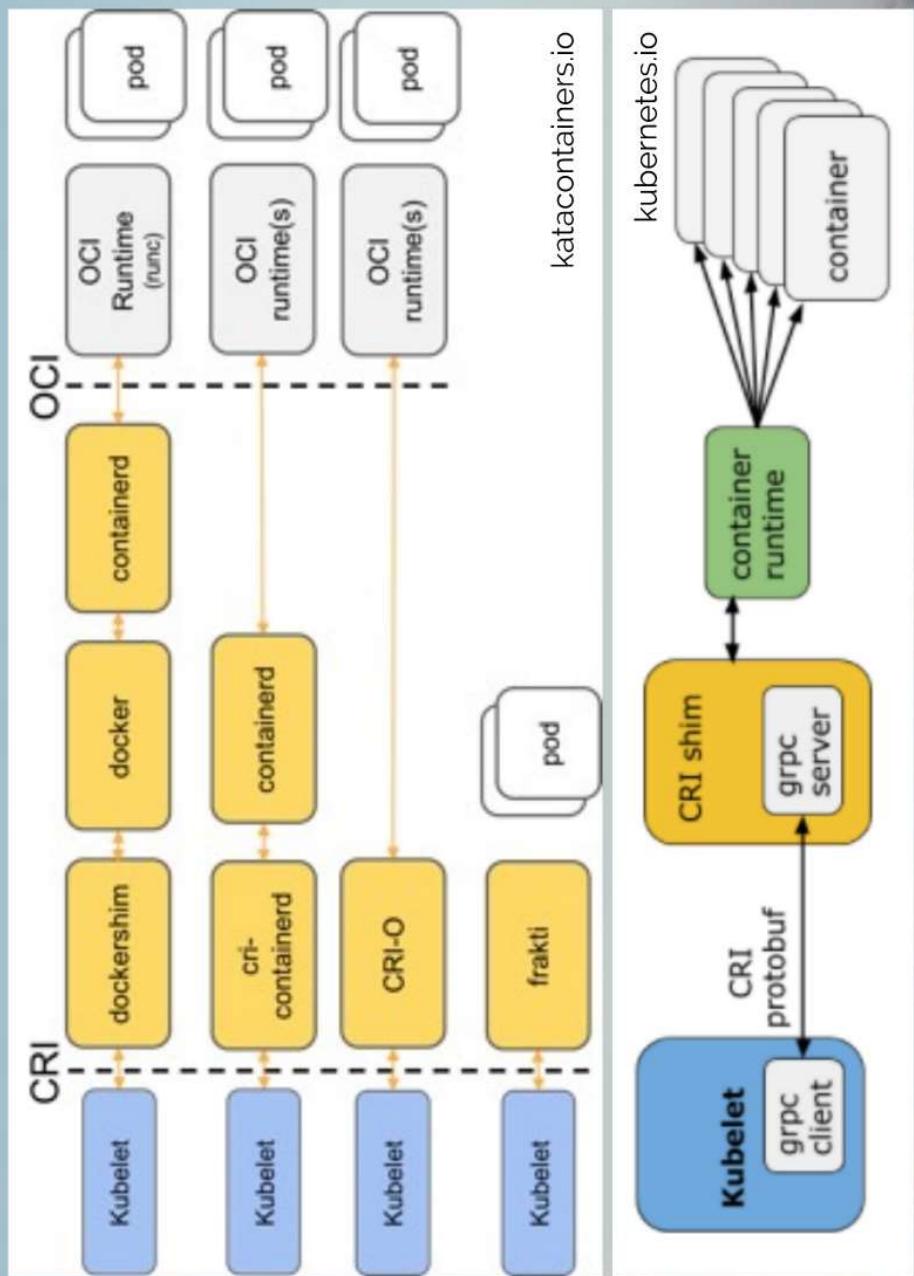
Tag

A label that uniquely identifies multiple container images in a single repository.

Base Image

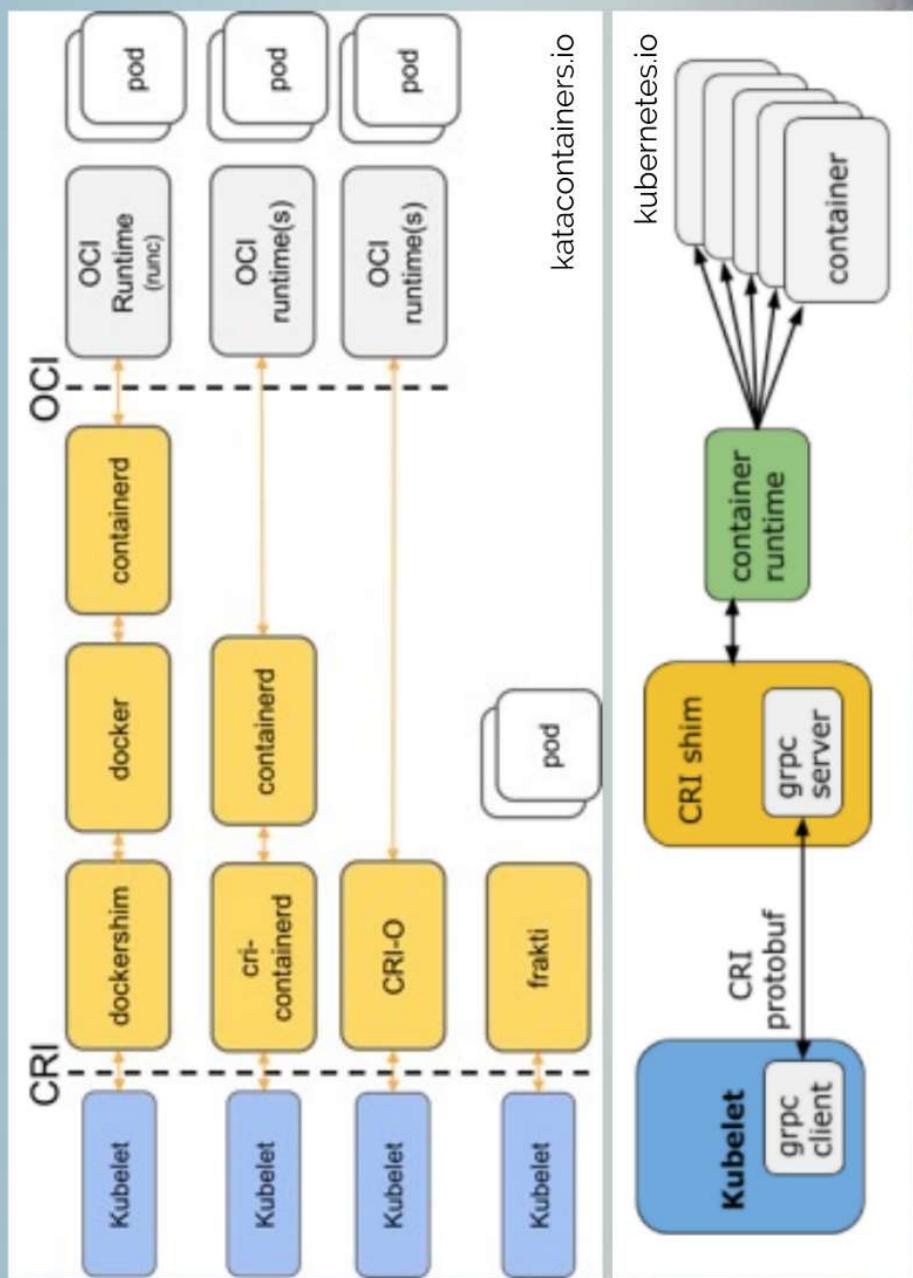
Most containers build upon base container images layers, and build more layers on top of them. Encourages reuse.

CRI and OCI



CRI and OCI

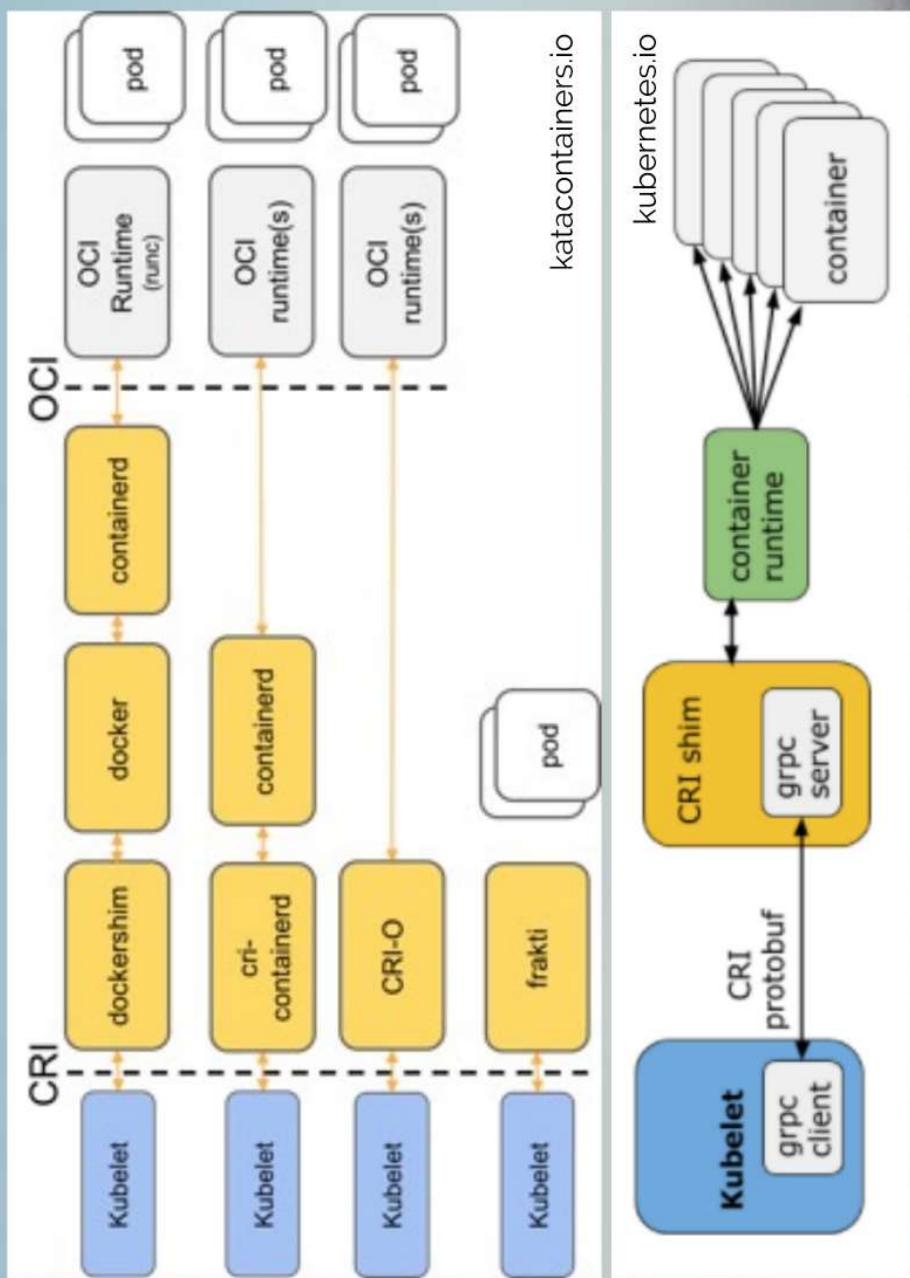
Container Runtime Interface



CRI and OCI

Container
Runtime Interface

Open Container
Initiative specification

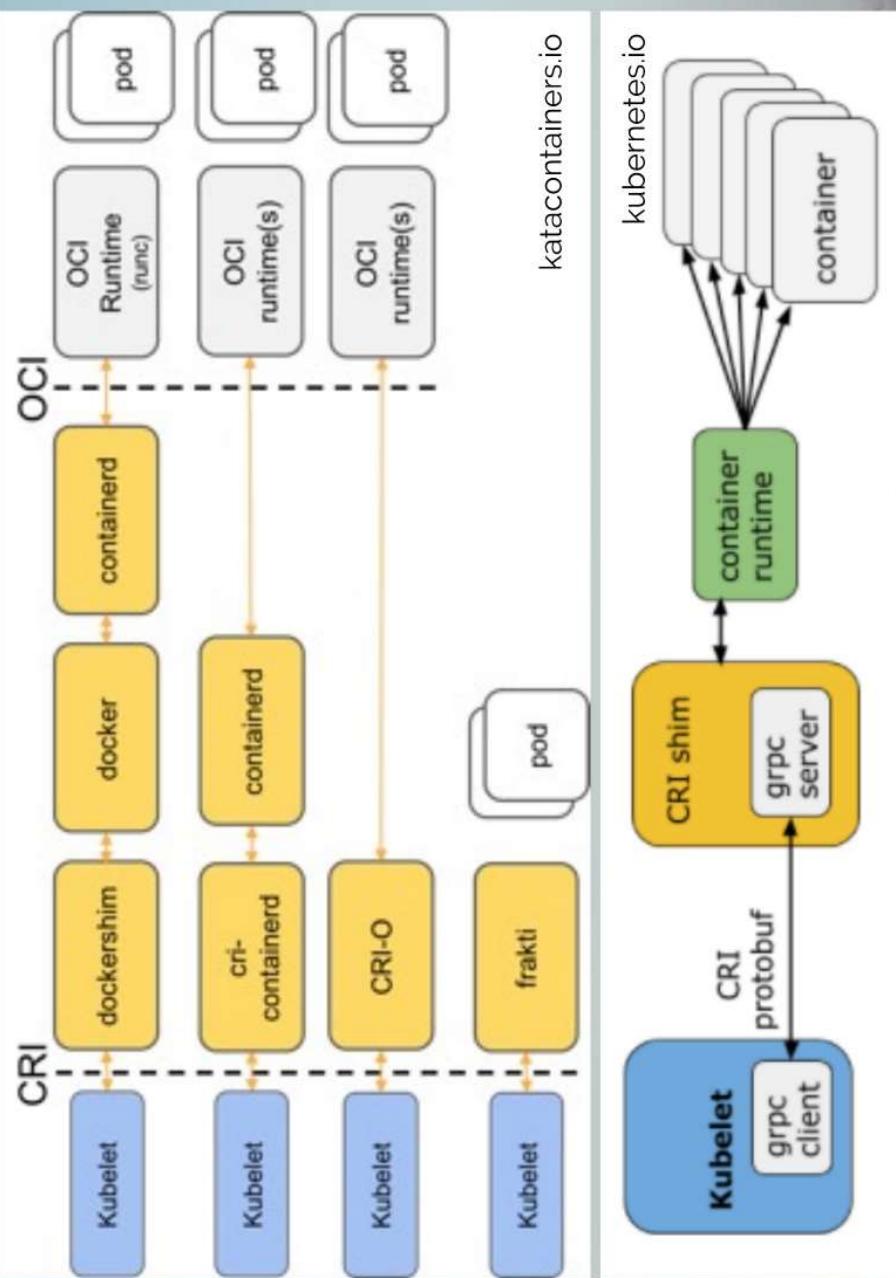


CONTAINER
INITIATIVE

CRI and OCI

Container Runtime Interface

Open Container Initiative specification



OCI container runtimes with Kubelet CRI

- Docker (nix Q4 2020)
- cri-o (OpenShift)
- **cri-containerd** (default)
- runc
- gVisor
- Kata containers
- Firecracker
- Nabla containers
- crun (c based)
- Rktlet (CRI for Rkt)
- rktnetes
- Frakti
- Singularity
- * Virtual Kubelet

CNCF Container Runtime Engines

 Running Apps

 Microservices

 APIs

 Polyglot

 Container Tools

 Developing Containers

 Rise of Containers

 Kubernetes Fundamentals

 Building Cloud Native Apps

 Pod Specifications

 Resources



Distillation Factors

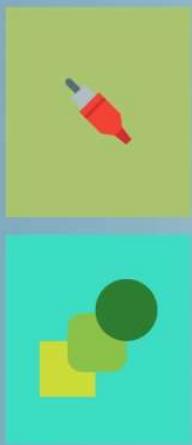
8 Factors

Distillation Factors

8 Factors

High cohesion

All features in the container are well defined, and used.
"Do one thing and do it well." -- Unix



Low coupling

Exposed interfaces are minimal, simple, and stable.



Distillation Factors

8 Factors

High cohesion

All features in the container are well defined, and used.
"Do one thing and do it well." -- Unix



Low coupling

Exposed interfaces are minimal, simple, and stable.

Immutable

Container not capable of or susceptible to change.

Idempotent

Multiple instances all behave the same, over time.

Distillation Factors

8 Factors

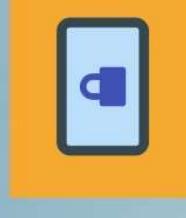
High cohesion

All features in the container are well defined, and used.
"Do one thing and do it well." -- Unix



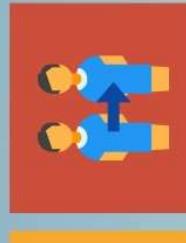
Low coupling

Exposed interfaces are minimal, simple, and stable.



Immutable

Container not capable of or susceptible to change.



Idempotent

Multiple instances all behave the same, over time.



Attack phobic

Security attack vectors are resisted as potential surfaces have been minimized.



Small images

Less is more. Reduce I/O pressure. Images are revised, stored, transported, downloaded and scaled.

Distillation Factors

8 Factors

High cohesion

All features in the container are well defined, and used.
"Do one thing and do it well." -- Unix

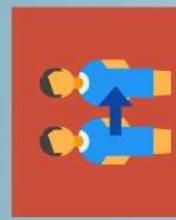
Low coupling

Exposed interfaces are minimal, simple, and stable.



Immutable

Container not capable of or susceptible to change.



Attack phobic

Security attack vectors are resisted as potential surfaces have been minimized.

Small images

Less is more. Reduce I/O pressure. Images are revised, stored, transported, downloaded and scaled.



Rapid startup

Containers are embraced as ephemeral, fragile, expendable, and scalable cattle.

Resource frugality

Every bit and CPU tick is a limited resource as containers are scaled throughout your datacenter.



Inspects, optimizes, and secures container images by removing unnecessary components

Java application images:
from ubuntu:14.04 - 743.6 MB → 100.3 MB

Node.js application images:

ubuntu:14.04 - 432MB → 14MB (by 30.85 x)
debian:jessie - 406MB → 25.1MB (16.21 x)
node:alpine - 66.7MB → 34.7MB (1.92 x)
node:distroless - 72.7MB → 39.7MB (1.83 x)

Python application images:

ubuntu:14.04 - 438MB → 16.8MB (by 25.99 x)
python:2.7-alpine - 84.3MB → 23.1MB (by 3.65 x)
python:2.7.15 - 916MB → 27.5MB (by 33.29 x)
centos:7 - 647MB → 23MB (by 28.57 x)
centos/python-27-centos7 - 700MB → 24MB (29.01 x)
python2.7:distroless - 60.7MB → 18.3MB (by 3.32 x)



Two Delicious Architectures

Containers

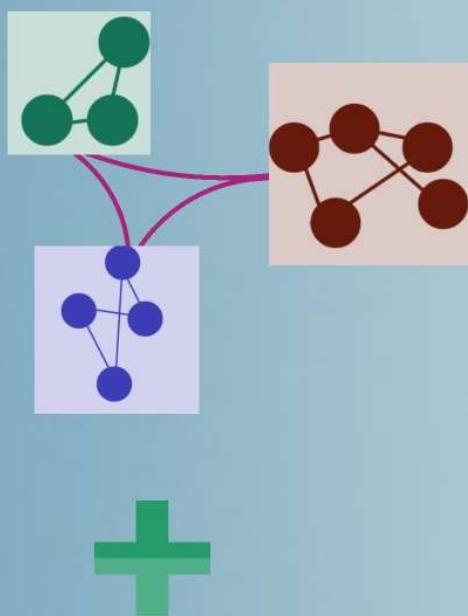


Two Delicious Architectures

Containers



Microservices

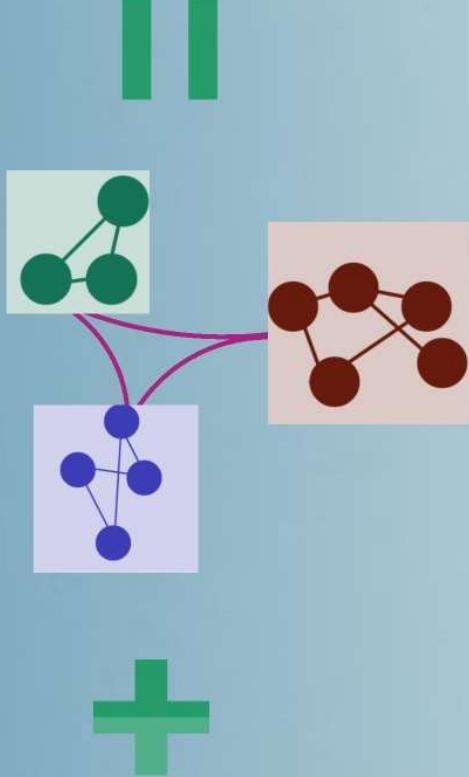


Two Delicious Architectures

Containers



Microservices



Tasty Combination



Microservices

- Apps as suites of independently deployable services
- Modular: High cohesion, low coupling
- Often own small datastore
- Scaling at a micro level
- Do one thing and do it well
- Sometimes they feel like OO classes
- Monolith antimatter

Microservices defined



James Lewis



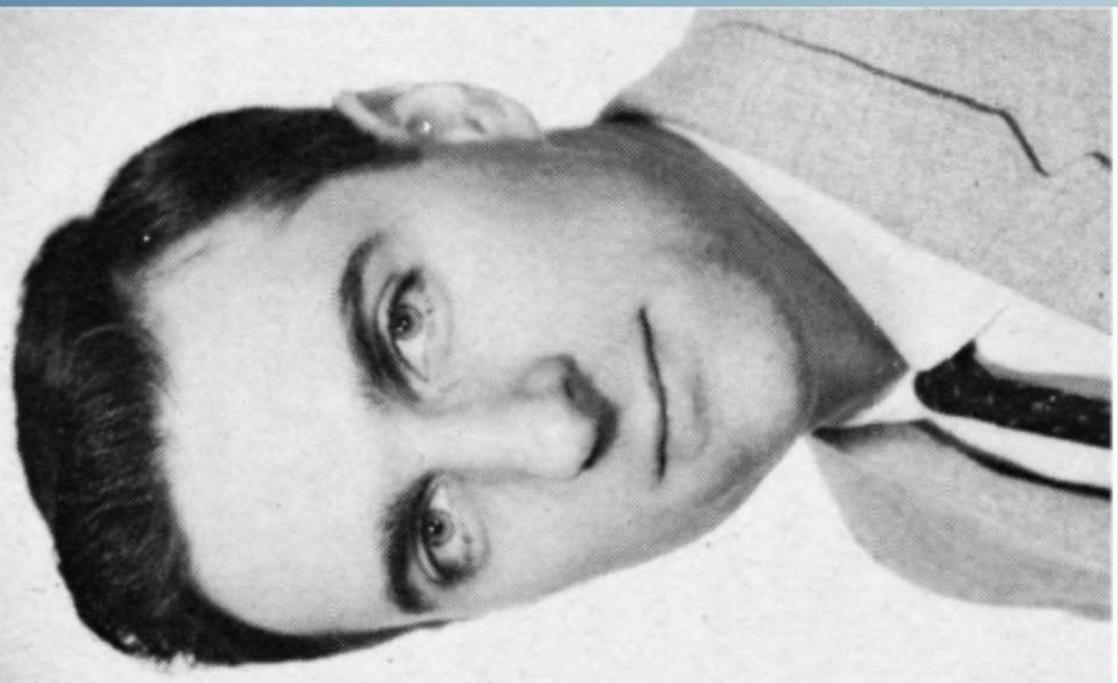
- Unix philosophy

Martin Fowler

Team organization

"Any organization that designs a system (defined more broadly here than just information systems) will inevitably produce a design whose structure is a copy of the organization's communication structure."

— Melvin E. Conway, 1968



Each team continuously delivers solutions, independently

Team

Communication and size

Jeff Bezos

- No matter how large your company gets, individual teams shouldn't be larger than what two pizzas can feed.

- API Mandate: All teams only communicate through APIs, else you are fired.

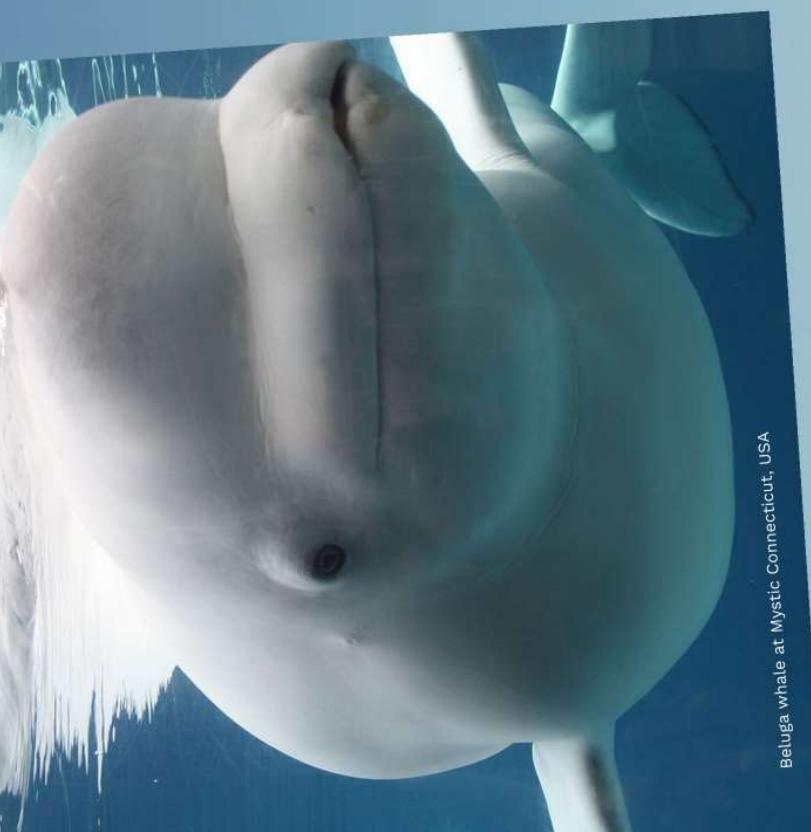


space-based architecture

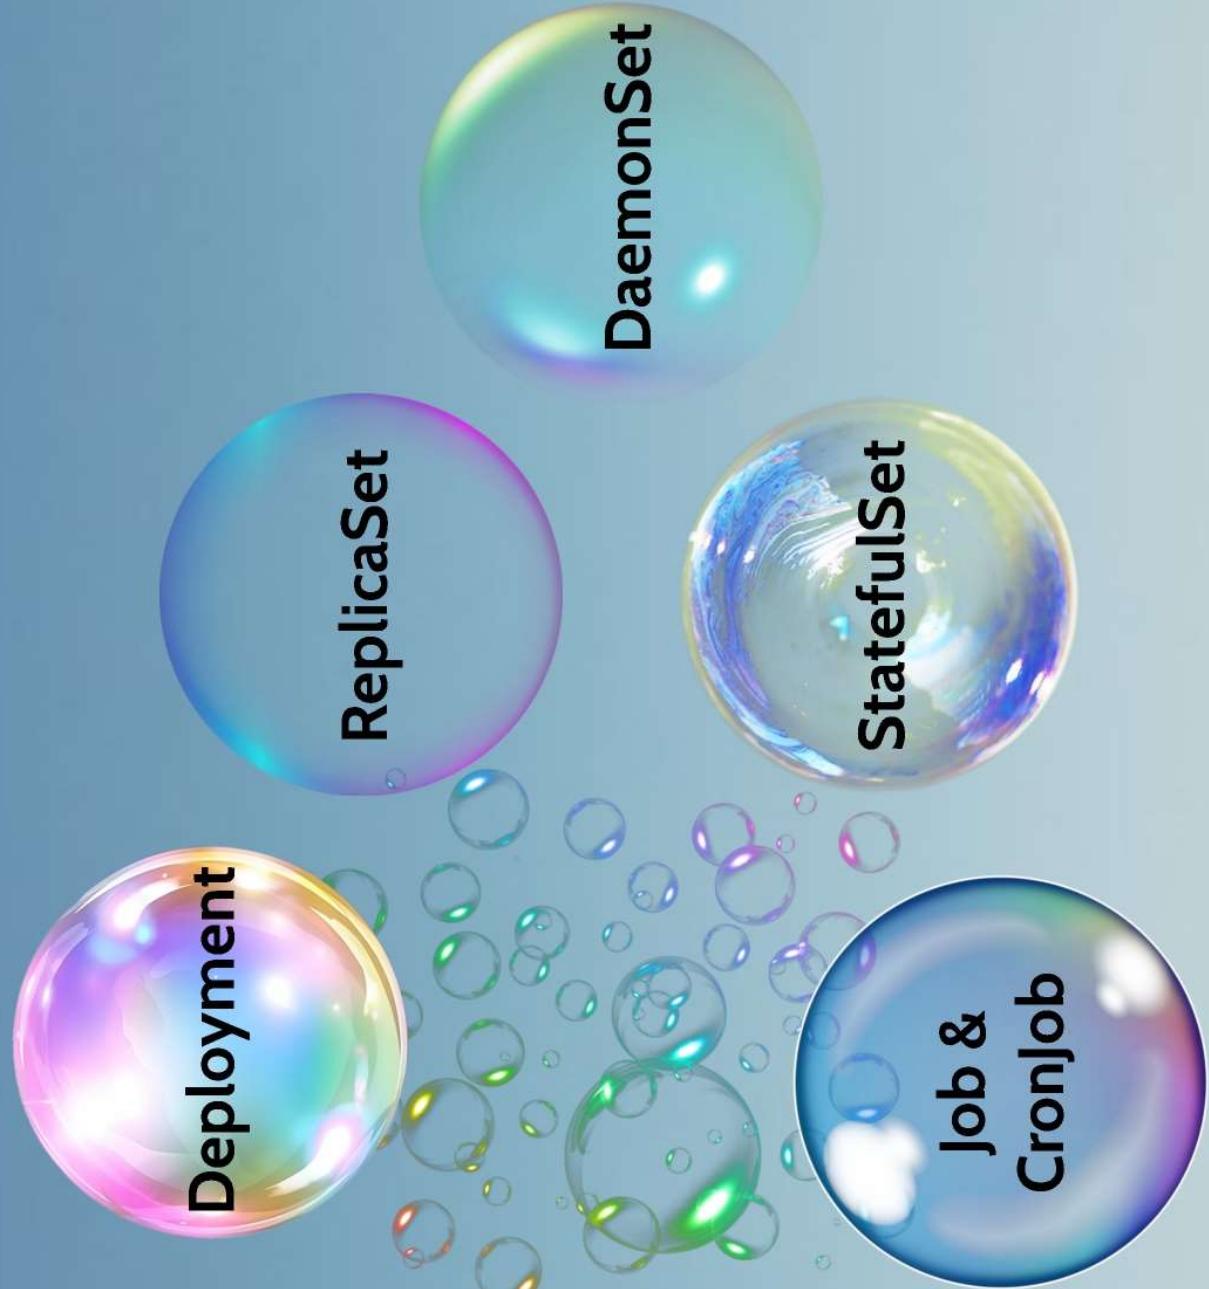
space-based architecture

	layered monolith	microkernel	microservices	service-based	event-driven	space-based
agility	★	★★	★★★	★★★★	★★★★★	★★★★★★
deployment	★	★★	★★★	★★★★	★★★★★	★★★★★★
testability	★★	★★★	★★★★	★★★★★	★★★★★★	★★★★★★★
performance	★★	★★★	★★★★	★★★★★	★★★★★★	★★★★★★★
scalability	★	★★	★★★	★★★★	★★★★★	★★★★★★
elasticity	★	★★	★★★	★★★★	★★★★★	★★★★★★
simplicity	★★	★★★	★★★★	★★★★★	★★★★★★	★★★★★★★
fault-tolerance	★	★★	★★★	★★★★	★★★★★	★★★★★★
evolvability	★	★★	★★★	★★★★	★★★★★	★★★★★★
total cost	★★	★★★	★★★★	★★★★★	★★★★★★	★★★★★★★

**Five kinds
Pods used by
built-in workload resources**



Beluga whale at Mystic, Connecticut, USA



kind: Deployment

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nginx-deployment
```

```
  labels:
```

```
    app: nginx
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
  matchLabels:
```

```
    app: nginx
```

```
  template:
```

```
    metadata:
```

```
    labels:
```

```
      app: nginx
```

```
spec:
```

```
  containers:
```

- name: nginx
- image: nginx:1.7.9
- ports:
- containerPort: 80



- Most common
- Rollout / Rollback
- Pause
- Image updates