

kind: ReplicaSet

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend
spec:
  replicas: 3
  selector:
  matchLabels:
    tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
  spec:
    containers:
      - name: php-redis
        image: frontend:v3
```



kind: ReplicaSet

- Define copies
- HPA Declares scaling
- Use Deployment instead



```
apiVersion: apps/v1
```

```
kind: ReplicaSet
```

```
metadata:
```

```
  name: frontend
```

```
  labels:
```

```
    app: guestbook
```

```
    tier: frontend
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
  matchLabels:
```

```
  tier: frontend
```

```
  template:
```

```
    metadata:
```

```
    labels:
```

```
    tier: frontend
```

```
  spec:
```

```
  containers:
```

```
    - name: php-redis
      image: frontend:v3
```

```
apiVersion: autoscaling/v1
```

```
kind: HorizontalPodAutoscaler
```

```
metadata:
```

```
  name: frontend-scaler
```

```
spec:
```

```
  scaleTargetRef:
```

```
    kind: ReplicaSet
```

```
    name: frontend
```

```
    minReplicas: 3
```

```
    maxReplicas: 10
```

```
    targetCPUUtilizationPercentage: 50
```

Declarative
Deployment

```
kubectl autoscale rs frontend --  
max=10 --min=3 --cpu-percent=50
```

Declarative
Horizontal

Horizontal Pod Autoscaler



SCENARIO

Kubernetes Observability: Scaling Your Applications, Automatically

Discover how the Horizontal Pod Autoscaler (HPA) can give you consistent performance while saving costs

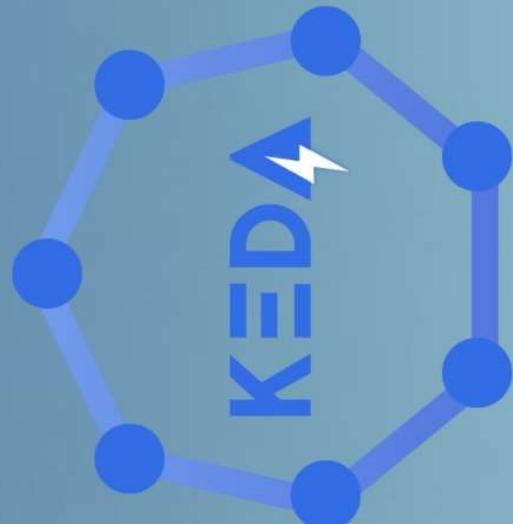
Lab

Explore

- Install metrics-server for gathering metrics,
- Install pod that can be scaled
- Define scaling rules and number of pods to scale
- Increase service demand to trigger scaling up
- Observe scaling up and down.

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: frontend-scaler
spec:
  scaleTargetRef:
    kind: ReplicaSet
    name: frontend
    minReplicas: 3
    maxReplicas: 10
    targetCPUUtilizationPercentage: 50
```

Operator for Event-driven Autoscaling

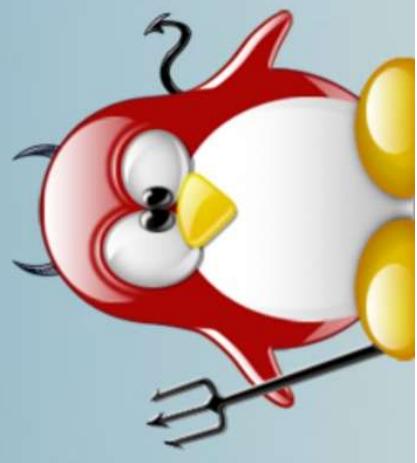
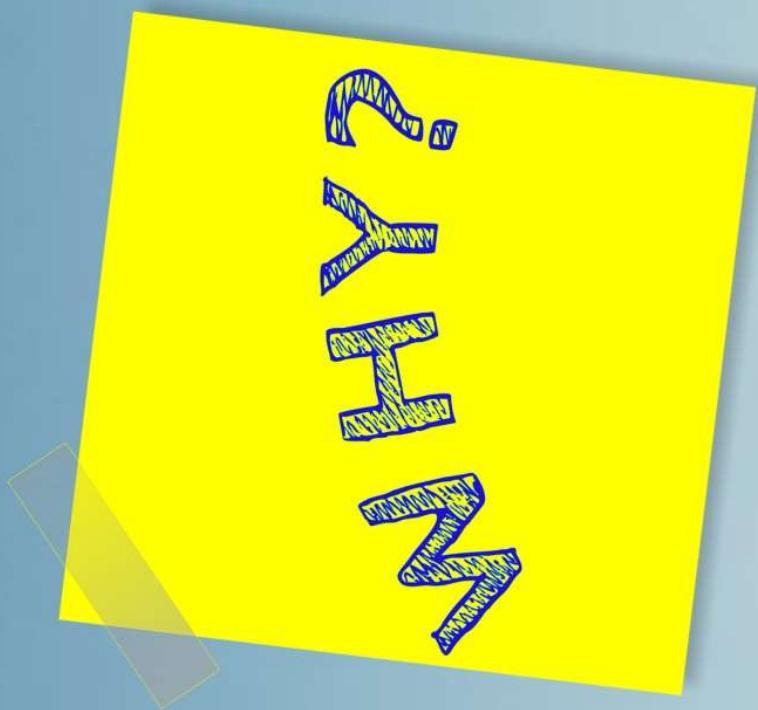


```
apiVersion: keda.k8s.io/v1alpha1
kind: ScaledObject
metadata:
  name: redis-scaledobject
  namespace: my-project
spec:
  scaleTargetRef:
    deploymentName: votes
  triggers:
  - type: redis
    metadata:
      address: REDIS_ADDRESS
  listName: myList
  listLength: "10"
  authenticationRef:
    name: trigger-auth-redis-secret
```

- Builds on HPA
- Context aware scaling
- Knative integration
- Scale to zero
- 64 scalers: CPU, memory, Metrics, etcd, MySQL, Datadog, Azure monitor, Azure Blob storage, ...

kind: DaemonSet

Pod on every Node



kind: DaemonSet

Pod on every Node

- Logging
- Metrics
- App and resource intelligence

- Taints and tolerations

```
apiVersion: apps/v1
```

```
kind: DaemonSet
```

```
metadata:
```

```
  name: fluentd-elasticsearch
```

```
labels:
```

```
  k8s-app: fluentd-logging
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      name: fluentd-elasticsearch
```

```
    template:
```

```
      metadata:
```

```
        labels:
```

```
          name: fluentd-elasticsearch
```

```
      spec:
```

```
        containers:
```

```
          - name: fluentd-elasticsearch
            image: fluentd:v2.5.2
            volumeMounts:
```

```
            - name: varlog
```

```
              mountPath: /var/log
```

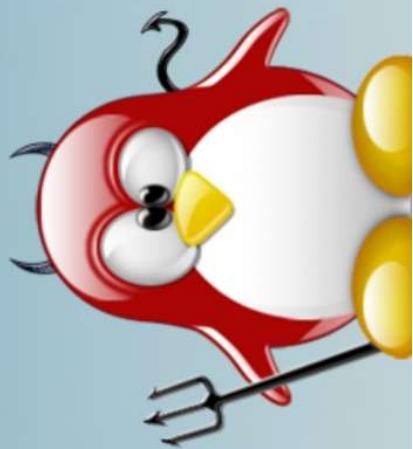
```
            - name: varlibdockercontainers
              mountPath: /var/lib/docker/containers
              readOnly: true
            terminationGracePeriodSeconds: 30
            volumes:
```

```
              - name: varlog
                hostPath:
```

```
                  path: /var/log
```

```
              - name: varlibdockercontainers
                hostPath:
```

```
                  path: /var/lib/docker/containers
```



kind: StatefulSet

- Sticky identity
- Connected volumes
- Distributed managers
- Consistent ID each pod
- Ordered members
- Each has ordinal index
- Ordered startup,
reverse shutdown



kind: StatefulSet

- Sticky identity
- Connected volumes
- Distributed managers
- Consistent ID each pod
- Ordered members
- Each has ordinal index
- Ordered startup, reverse shutdown

Examples



Datastores



kind: StatefulSet

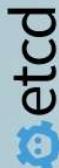
- Sticky identity
- Connected volumes
- Distributed managers
- Consistent ID each pod
- Ordered members
- Each has ordinal index
- Ordered startup, reverse shutdown

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: k8s.gcr.io/nginx-slim:0.8
  ports:
    - containerPort: 80
      name: web
  volumeMounts:
    - name: www
      mountPath: /usr/share/nginx/html
  volumeClaimTemplates:
    - metadata:
        name: www
      spec:
        accessModes: [ "ReadWriteOnce" ]
        resources:
          requests:
            storage: 1Gi
```

Examples



Datastores



kind: Job

```
apiVersion: batch/v1  
kind: Job
```

```
metadata:
```

```
  name: pi-with-timeout
```

```
spec:
```

```
  backoffLimit: 5
```

```
  activeDeadlineSeconds: 100
```

```
  template:
```

```
    spec:
```

```
      restartPolicy: Never
```

```
      containers:
```

```
        - name: pi
```

```
          image: perl
```

```
          command: ["perl", "-Mbignum=bpi", "-wle", "print  
bpi(2000)"]
```

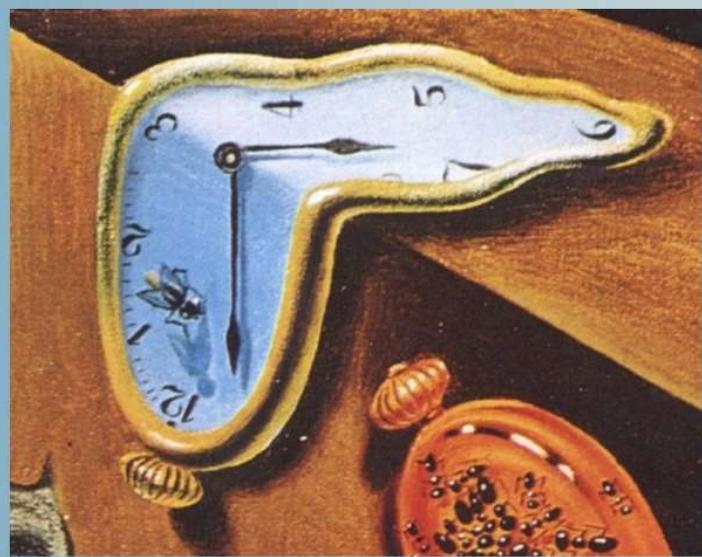
- Triggered from events
 - Run in Pods then stop
 - Parallel or not
 - Failure policies
- Use patterns
 - Installations
 - Cleanups
 - Builds
- Datastore updates
 - Helm, Terraform,
 - Liquibase
- Provisioning
 - Processing



kind: CronJob

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: hello
spec:
  schedule: "* */1 * * *"
  jobTemplate:
    spec:
      template:
        spec:
```

- Reoccurring work
- Periodic
- Trigger based on time



 Running Apps

 Microservices

 APIs

 Polyglot

 Container Tools

 Developing Containers

 Rise of Containers

Kubernetes Fundamentals



Building Cloud Native Apps

 Pod Specifications

 Resources

 Pod Workloads



Fin

Podulation



kubernetes

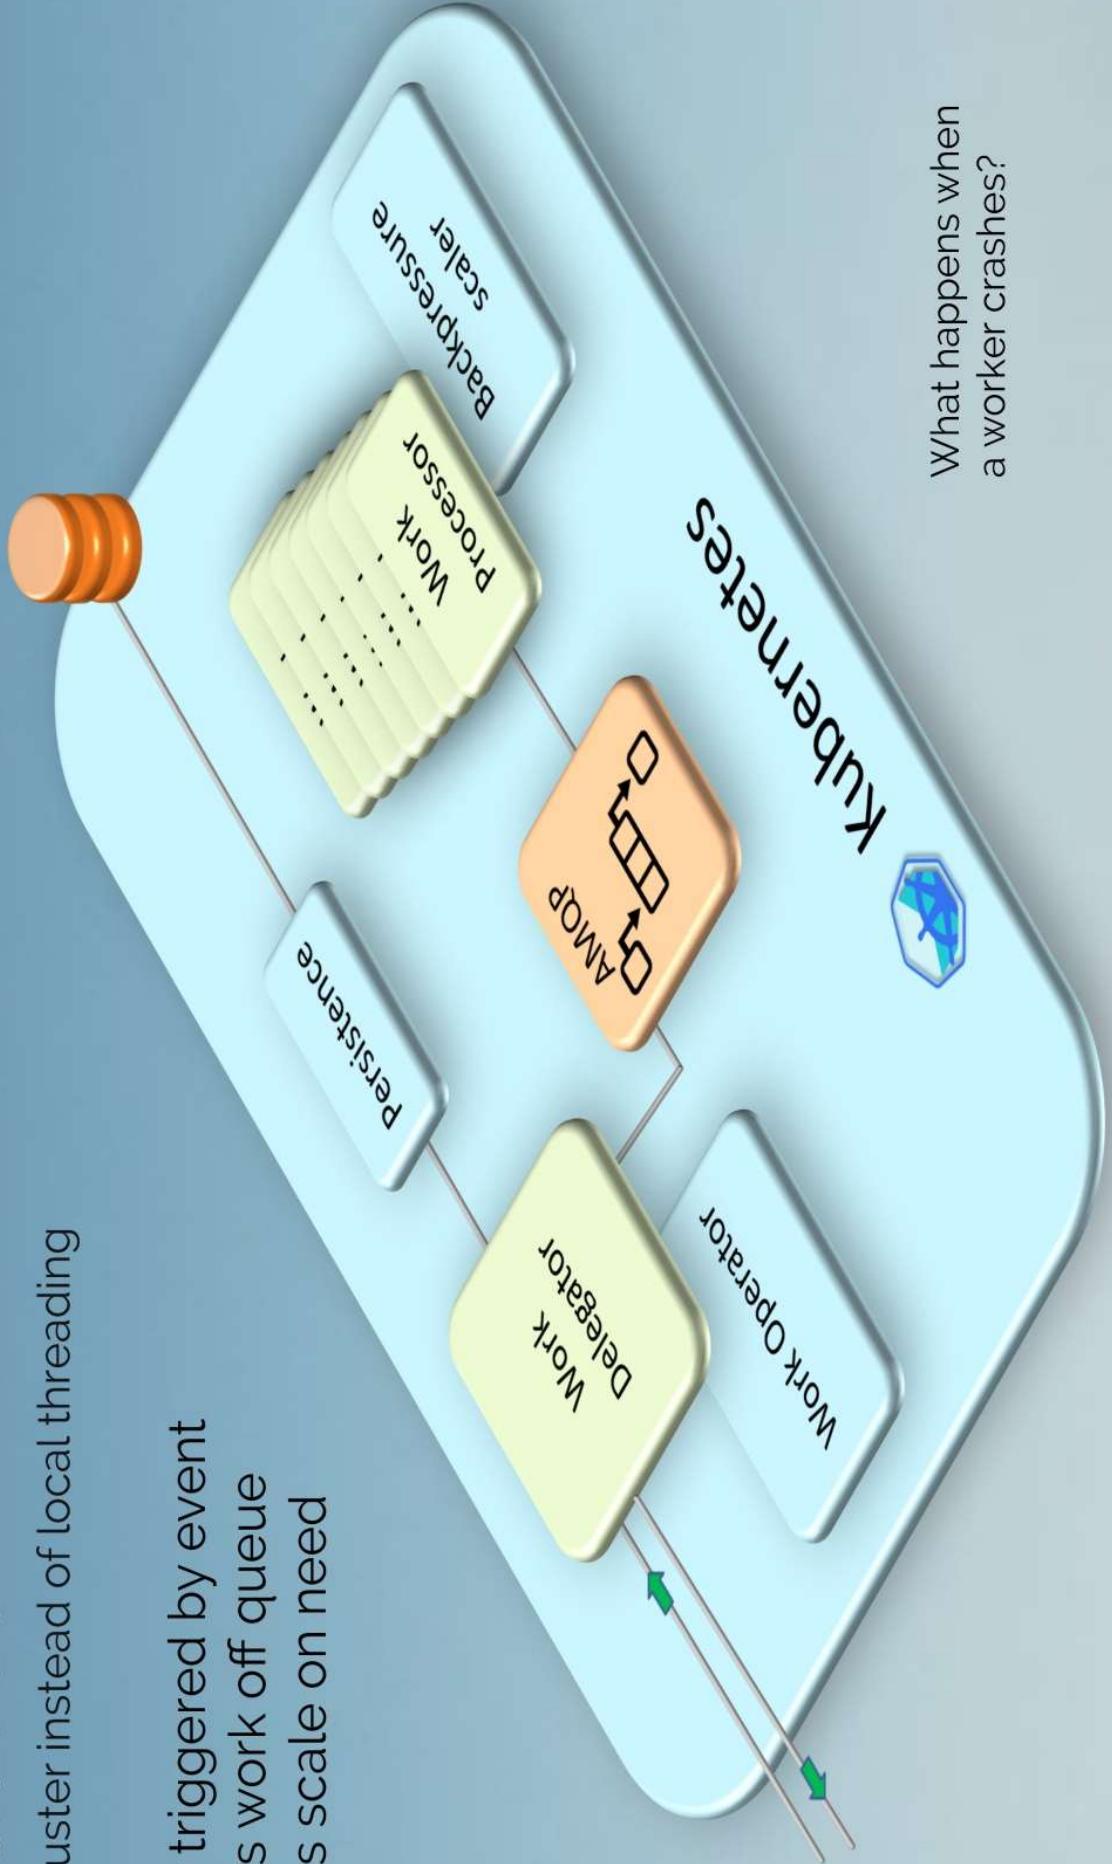
Software Architecture

Patterns and practices
of grouping collections of applications
distributed across a cluster.

Work Queue

Scale with cluster instead of local threading

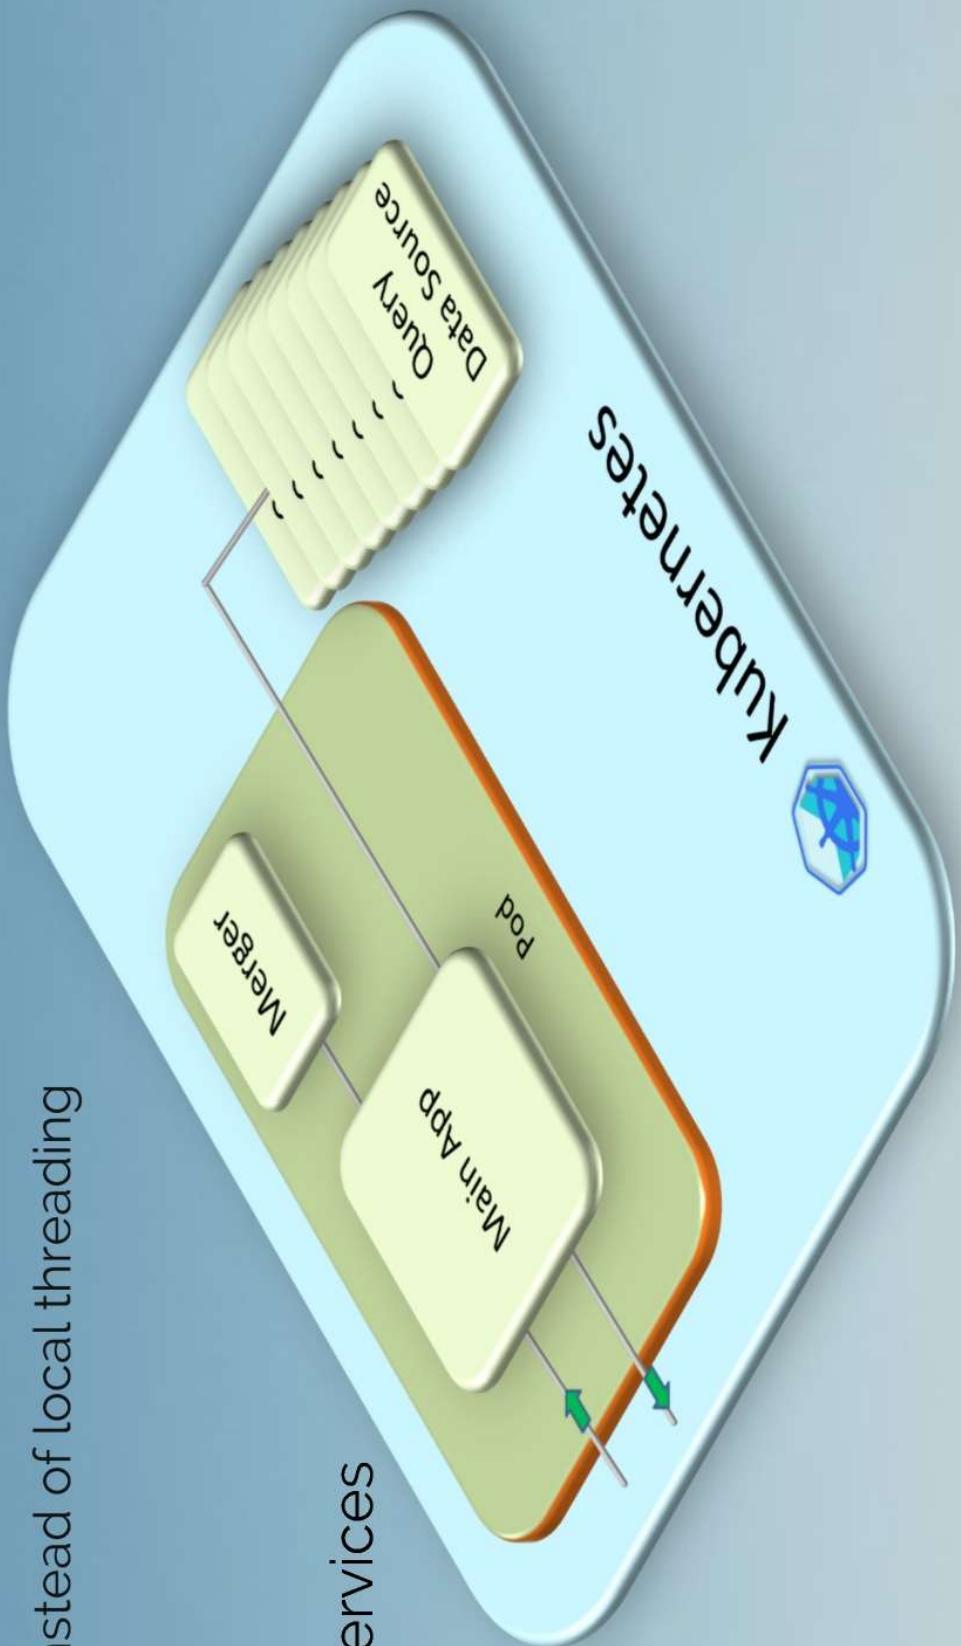
- Delegator triggered by event
- Processors work off queue
- Processors scale on need



Scatter and Gather

Scale with cluster instead of local threading

- Fan out
- Query multiple services
- Gather data
- Merge results
- Same pod or not



Grouped Containers in Pods

- Pods are atomic units
- Shared Linux namespace
- Assigned IP addresses 10.X.X.X
- localhost between containers
- Shared mounts
- Shared port space



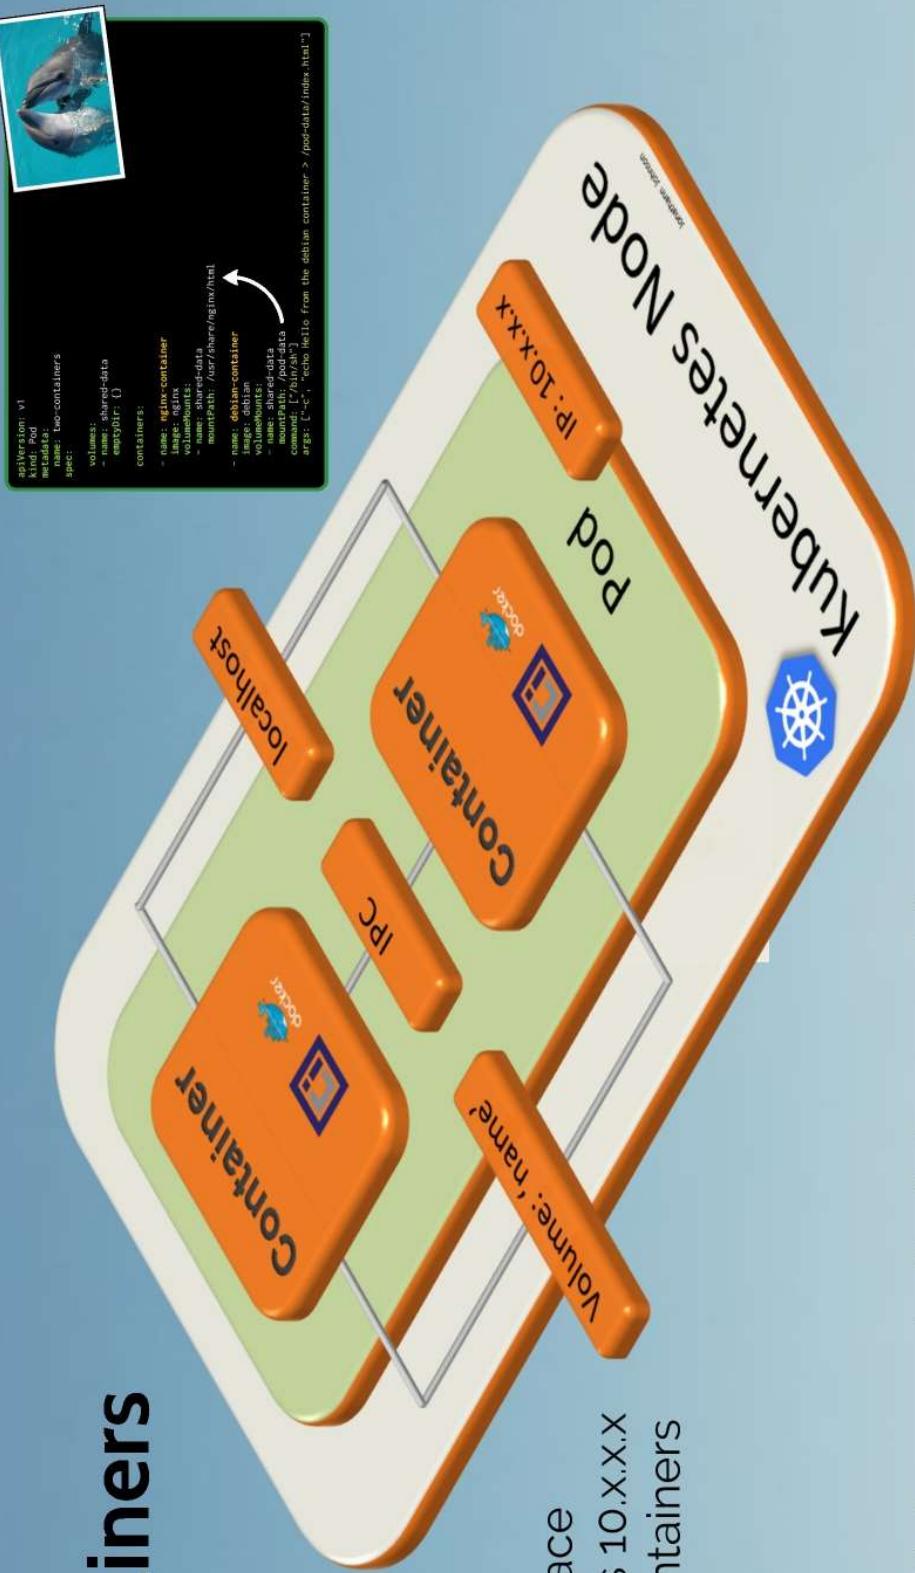
Fast communication between containers

- Shared PersistenceVolume mount: "emptyDir"
- localhost
- Linux inter-process communication (IPC)
- Signals, pipes, sockets, message queues, semaphores, shared memory



Grouped Containers in Pods

- Pods are atomic units
- Shared Linux namespace
- Assigned IP addresses 10.X.X.X
- localhost between containers
- Shared mounts
- Shared port space



Fast communication between containers

- Shared PersistenceVolume mount: "emptyDir"
- localhost
- Linux inter-process communication (IPC)
- Signals, pipes, sockets, message queues, semaphores, shared memory



```
apiVersion: v1
kind: Pod
metadata:
  name: two-containers
spec:
  volumes:
    - name: shared-data
      emptyDir: {}
  containers:
    - name: nginx-container
      image: nginx
      volumeMounts:
        - name: shared-data
          mountPath: /usr/share/nginx/html
    - name: debian-container
      image: debian
      volumeMounts:
        - name: shared-data
          mountPath: /pod-data
          command: ["/bin/sh"]
          args: ["-c", "echo Hello from the debian container > /pod-data/index.html"]
```

 Running Apps

 Microservices

 Pod Workloads

 APIs

 Polyglot

 Container Tools

 Developing Containers

 Rise of Containers

 Building Cloud Native Apps

Kubernetes Fundamentals



 Pod Specifications

 Resources

 Learning Channels



Pod Specifications

Metadata

InitContainers

Images

Sidecar

Context

Affinity

Volumes

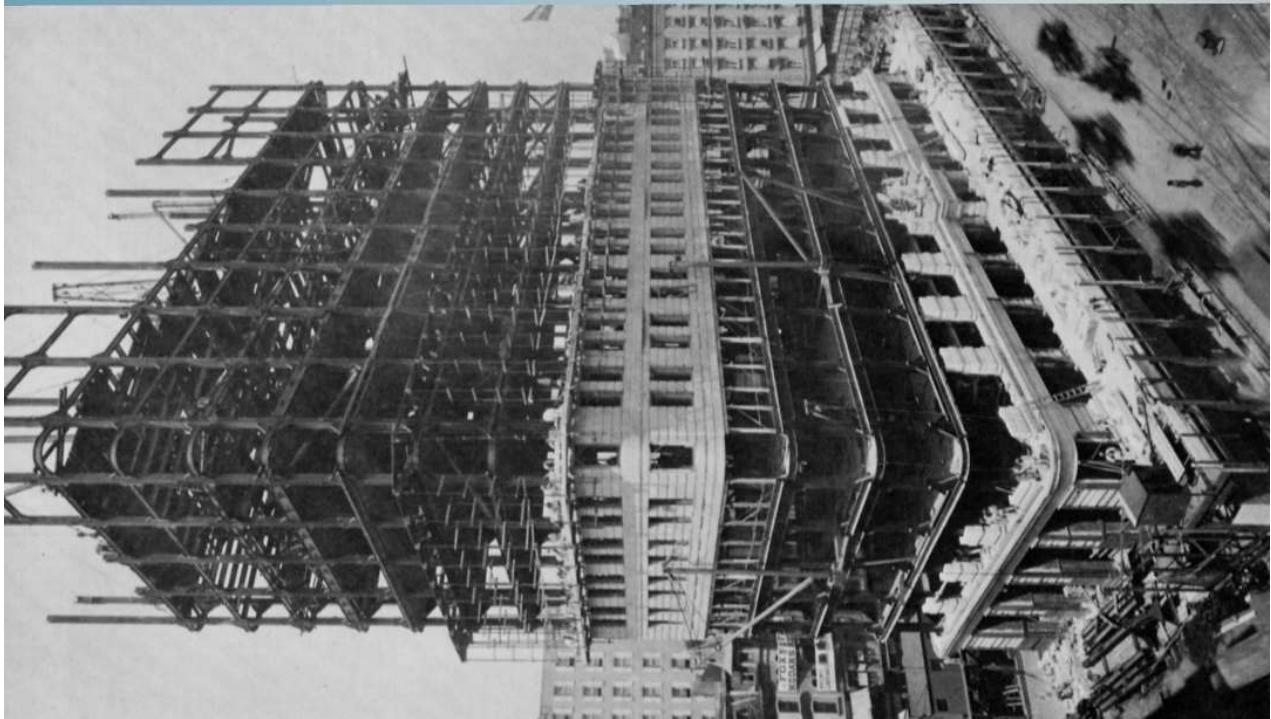
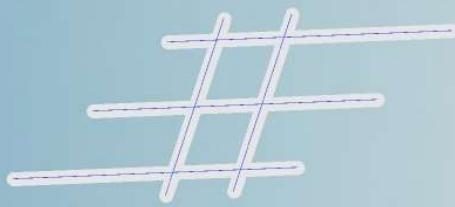
Disruptions

Probes

Ephemeral
Containers

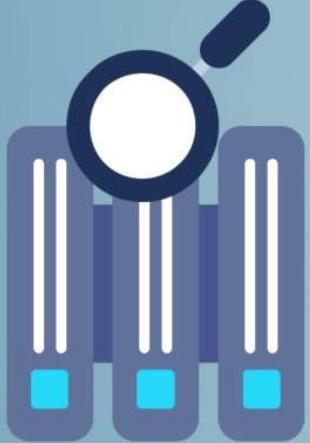
Quotas

Security



Metadata

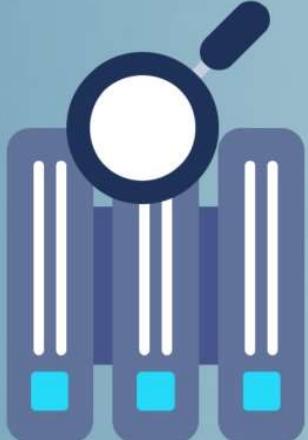
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-two-deployment
  annotations:
    devStage: beta
  labels:
    type: webserver
spec:
  selector:
    matchLabels:
      app: nginx-two
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx-two
spec:
  containers:
    - name: nginx-two
      image: nginx:1.25-alpine
      ports:
        - containerPort: 80
```



Metadata

Name: unique for
resource in namespace

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-two-deployment
  annotations:
    devStage: beta
  labels:
    type: webserver
spec:
  selector:
    matchLabels:
      app: nginx-two
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx-two
    spec:
      containers:
        - name: nginx-two
          image: nginx:1.25-alpine
          ports:
            - containerPort: 80
```

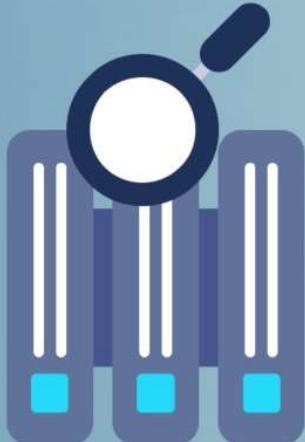


Name: unique for
container in pod

Metadata

Name: unique for resource in namespace

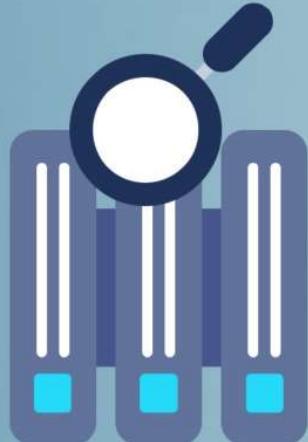
Annotate: Attach arbitrary non-identifying metadata to objects



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-two-deployment
  annotations:
    devStage: beta
  labels:
    type: webserver
spec:
  selector:
    matchLabels:
      app: nginx-two
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx-two
    spec:
      containers:
        - name: nginx-two
          image: nginx:1.25-alpine
          ports:
            - containerPort: 80
```

Name: unique for container in pod

Metadata



Name: unique for resource in namespace

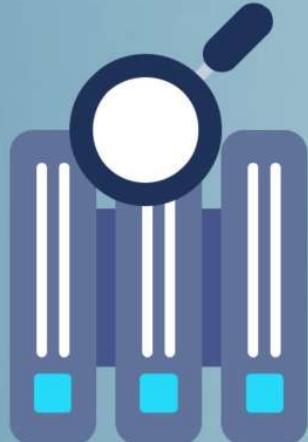
Annotate: Attach arbitrary non-identifying metadata to objects

Labels: For categorizing to match, search, and filter objects

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-two-deployment
  annotations:
    devStage: beta
  labels:
    type: webserver
spec:
  selector:
    matchLabels:
      app: nginx-two
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx-two
spec:
  containers:
    - name: nginx-two
      image: nginx:1.25-alpine
      ports:
        - containerPort: 80
```

Name: unique for container in pod

Metadata



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-two-deployment
  annotations:
    devStage: beta
  labels:
    type: webserver
spec:
  selector:
    matchLabels:
      app: nginx-two
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx-two
    spec:
      containers:
        - name: nginx-two
          image: nginx:1.25-alpine
          ports:
            - containerPort: 80
```

Name: unique for resource in namespace

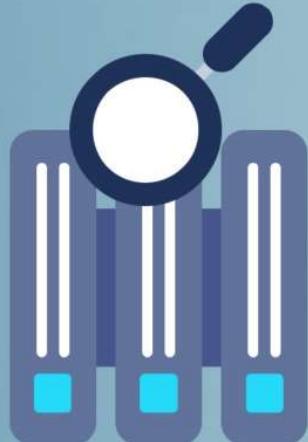
Annotate: Attach arbitrary non-identifying metadata to objects

Labels: For categorizing to match, search, and filter objects

Namespace: Not often applied in manifest

Name: unique for container in pod

Metadata



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-two-deployment
  annotations:
    devStage: beta
  labels:
    type: webserver
spec:
  selector:
    matchLabels:
      app: nginx-two
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx-two
    spec:
      containers:
        - name: nginx-two
          image: nginx:1.25-alpine
          ports:
            - containerPort: 80
```

Name: unique for resource in namespace

Annotate: Attach arbitrary non-identifying metadata to objects

Labels: For categorizing to match, search, and filter objects

Namespace: Not often applied in manifest

💡 Q: Label vs Annotation?
A: Use annotation when you don't need to search on a label

Name: unique for container in pod

Community defines Recommended Labels
• Well-Known Labels, Annotations and Taints

- **name:** nginx-two
image: nginx:1.25-alpine
ports:
- containerPort: 80

Images

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
      replicas: 3
      template:
        metadata:
          labels:
            app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx:1.19.3-alpine
        ports:
          - name: http-port
            containerPort: 80
        imagePullSecrets:
          - name: regcred
```

1 or more containers per pod

Image pulled from repos...

- Reference image tag:version
- imagePullPolicy
- imagePullSecrets (ifNotPresent, Always, Never)

Each Container Has...

- Name
- Ports
- Probes
- Volumes
- Environment vars
- CPU / memory quotas
- Hooks: PostStart & PreStop

THE TWELVE-FACTOR APP

Factor III: Config

Store config in the environment

An app's config is everything that is likely to vary between deploys (staging, production, developer environments, etc). This includes:

- Resource handles to the databaseMemcached, and other backing services
- Credentials to external services such as Amazon S3 or Twitter
- Per-deploy values such as the canonical hostname for the deploy

Where to store the environment var for the URL for a microservice connecting to MySQL?

```
MySQL_URL = jdbc:mysql://host1:33060/sakila
```

Where to store a password to access MySQL?

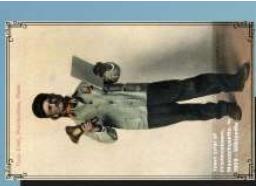
```
MySQL_PASSWORD = "shhhhh"
```

Declaring ConfigMaps and Secrets

Data can be linked to:

- Environment variables
- Command line parameters
- File mounts (read-only)

```
apiVersion: v1
kind: Pod
metadata:
  name: passable
spec:
  containers:
    - name: question-app
      image: grail-seeker
  env:
    - name: swallow-bird-type
      value: "African"
    - name: QUESTION_ONE
      valueFrom:
        configMapKeyRef:
          name: questions
          key: q1
    - name: QUESTION_TWO
      valueFrom:
        configMapKeyRef:
          name: questions
          key: q2
```



Declaring ConfigMaps and Secrets

Data can be linked to:

- Environment variables
- Command line parameters
- File mounts (read-only)

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: questions
data:
  q1: "What is your quest?"
  q2: "What is your name?"
```

```
apiVersion: v1
kind: Pod
metadata:
  name: passable
```

```
spec:
  containers:
    - name: question-app
      image: grail-seeker
  env:
    - name: swallow-bird-type
      value: "African"
    - name: QUESTION_ONE
      valueFrom:
        configMapKeyRef:
          name: questions
          key: q1
    - name: QUESTION_TWO
      valueFrom:
        configMapKeyRef:
          name: questions
          key: q2
```



Declaring ConfigMaps and Secrets

Data can be linked to:

- Environment variables
- Command line parameters
- File mounts (read-only)

```
apiVersion: v1
kind: Pod
metadata:
  name: passable
spec:
  containers:
    - name: question-app
      image: grail-seeker
  env:
    - name: swallow-bird-type
      value: "African"
    - name: QUESTION_ONE
      valueFrom:
        configMapKeyRef:
          name: questions
          key: q1
    - name: QUESTION_TWO
      valueFrom:
        configMapKeyRef:
          name: questions
          key: q2
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: questions
data:
  q1: "What is your quest?"
  q2: "What is your name?"
```



Declaring ConfigMaps and Secrets

Data can be linked to:

- Environment variables
- Command line parameters
- File mounts (read-only)

```
apiVersion: v1
kind: Pod
metadata:
  name: passable
spec:
  containers:
    - name: question-app
      image: grail-seeker
  env:
    - name: swallow-bird-type
      value: "African"
    - name: QUESTION_ONE
      valueFrom:
        configMapKeyRef:
          name: questions
          key: q1
    - name: QUESTION_TWO
      valueFrom:
        configMapKeyRef:
          name: questions
          key: q2
```



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: questions
data:
  q1: "What is your quest?"
  q2: "What is your name?"
```

Context



SCENARIO

Kubernetes Fundamentals: ConfigMaps and Secrets

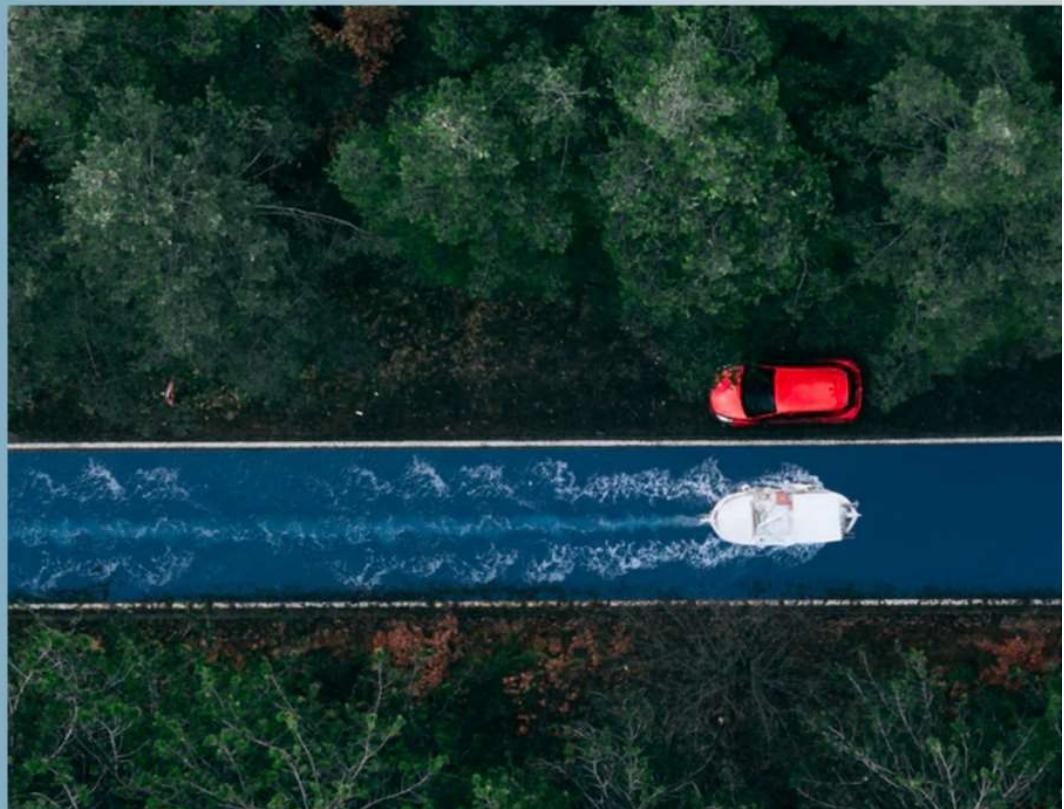
The fundamentals of ConfigMaps and Secrets

Lab



Explore

- Create configuration data
- Share context across deployments
- 3 techniques Pods access configuration data
- Keeping data outside of code



Access Secrets in Store



CONTAINER
STORAGE
INTERFACE

Secrets Store CSI Driver

The Secrets Store CSI Driver secrets-store.csi.k8s.io allows Kubernetes to mount multiple secrets, keys, and certs stored in enterprise-grade external secrets stores into their pods as a volume.

Several providers

- Akeyless Provider
- AWS Provider
- Azure Provider
- GCP Provider
- Vault Provider

Pod accesses K8s Secret via a file mount.
Auto rotation can set set for Secrets -- should be tested

```
apiVersion: secrets-store.csi.x-k8s.io/v1
kind: SecretProviderClass
metadata:
  name: my-provider
spec:
  provider: azure
  parameters:
```



Secrets with Azure



Set up Azure Keyvault with secrets (az commands)

Use the Azure Key Vault provider for Secrets Store CSI Driver in an AKS



Helm: Install Secret Store CIS Driver and Azure Keyvault Provider

- [ArtifactHub](#)
- [Homepage](#)

Secrets with Azure

Pod accesses Secret via a file mount

```
apiVersion: secrets-store.csi.x-k8s.io/v1
kind: SecretProviderClass
metadata:
  name: my-provider
spec:
  provider: azure
  parameters:
    usePodIdentity: "false"
    useManagedIdentity: "false"
    keyVaultName: "$KEYVAULT_NAME"
  objects: |
    array:
      - |
        objectName: $SECRET_NAME
        objectType: secret
        objectVersion: $SECRET_VERSION
      - |
        objectName: $KEY_NAME
        objectType: key
        objectVersion: $KEY_VERSION
        tenantId: "$TENANT_ID"
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: test-csi
```

```
spec:
```

```
  containers:
```

```
    - image: alpine
      name: alpine
      command:
```

```
      - "sh"
```

```
      - "-c"
```

```
      - "echo going to sleep... && sleep 10000"
```

```
volumeMounts:
```

```
  - name: secrets-store-inline
    mountPath: "/csi-secrets"
    readOnly: true
```

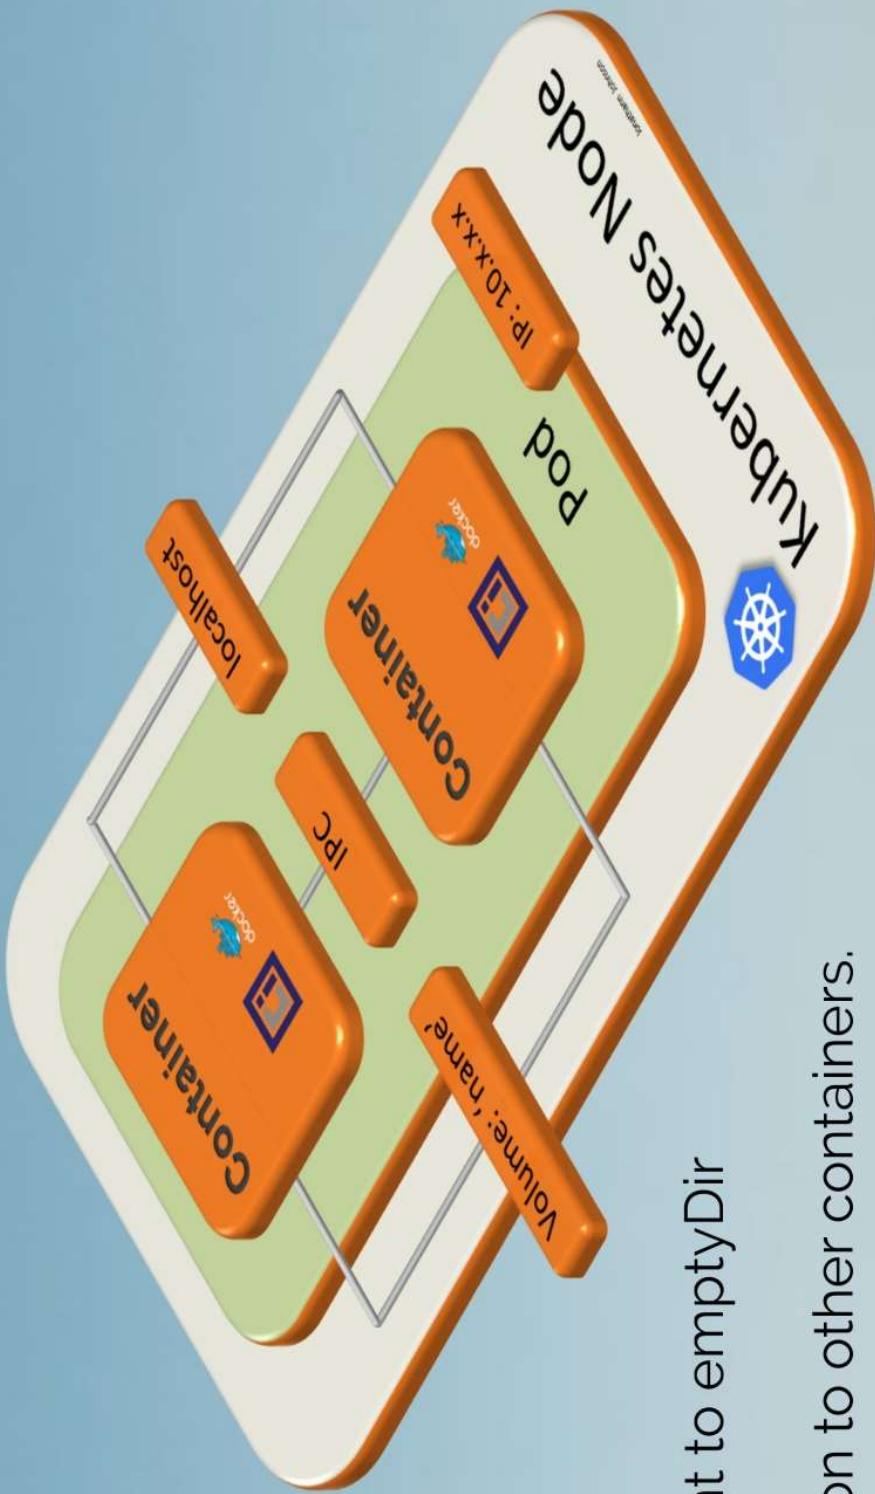
```
volumes:
```

```
  - name: secrets-store-inline
    csi:
      driver: secrets-store.csi.k8s.io
      readOnly: true
      volumeAttributes:
        secretProviderClass: "my-provider"
```



Volumes

Local Pod volume



Containers can mount to emptyDir
Use as cache
Use as communication to other containers.

Volumes

Local Node volume

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pd
spec:
  containers:
    - image: k8s.gcr.io/test-webserver
      name: test-container
  volumeMounts:
    - mountPath: /test-pd
      name: test-volume
```

Sometime containers need
access to Host Volume with
hostPath

hostPath mounts can typed as:

- DirectoryOrCreate
- **Directory**
- FileOrCreate
- File
- CharDevice
- BlockDevice
- (empty, no check)

```
hostPath:
# directory location on host
path: /data
# this field is optional
type: Directory
```

Volumes

Persistent volume

```
apiVersion: v1
kind: Pod
metadata:
  name: www
  labels:
    name: www
spec:
  containers:
    - name: www
      image: nginx:alpine
      ports:
        - containerPort: 80
          name: www
  volumeMounts:
    - name: www-persistent-storage
      mountPath: /usr/share/nginx/html
volumes:
  - name: www-persistent-storage
    persistentVolumeClaim:
      claimName: claim-http
```

Volumes

Persistent volume

```
apiVersion: v1
kind: Pod
metadata:
  name: www
  labels:
    name: www
spec:
  containers:
    - name: www
      image: nginx:alpine
      ports:
        - containerPort: 80
          name: www
  volumeMounts:
    - name: www-persistent-storage
      mountPath: /usr/share/nginx/html
volumes:
  - name: www-persistent-storage
    persistentVolumeClaim:
      claimName: claim-http
```

Volumes

Persistent volume

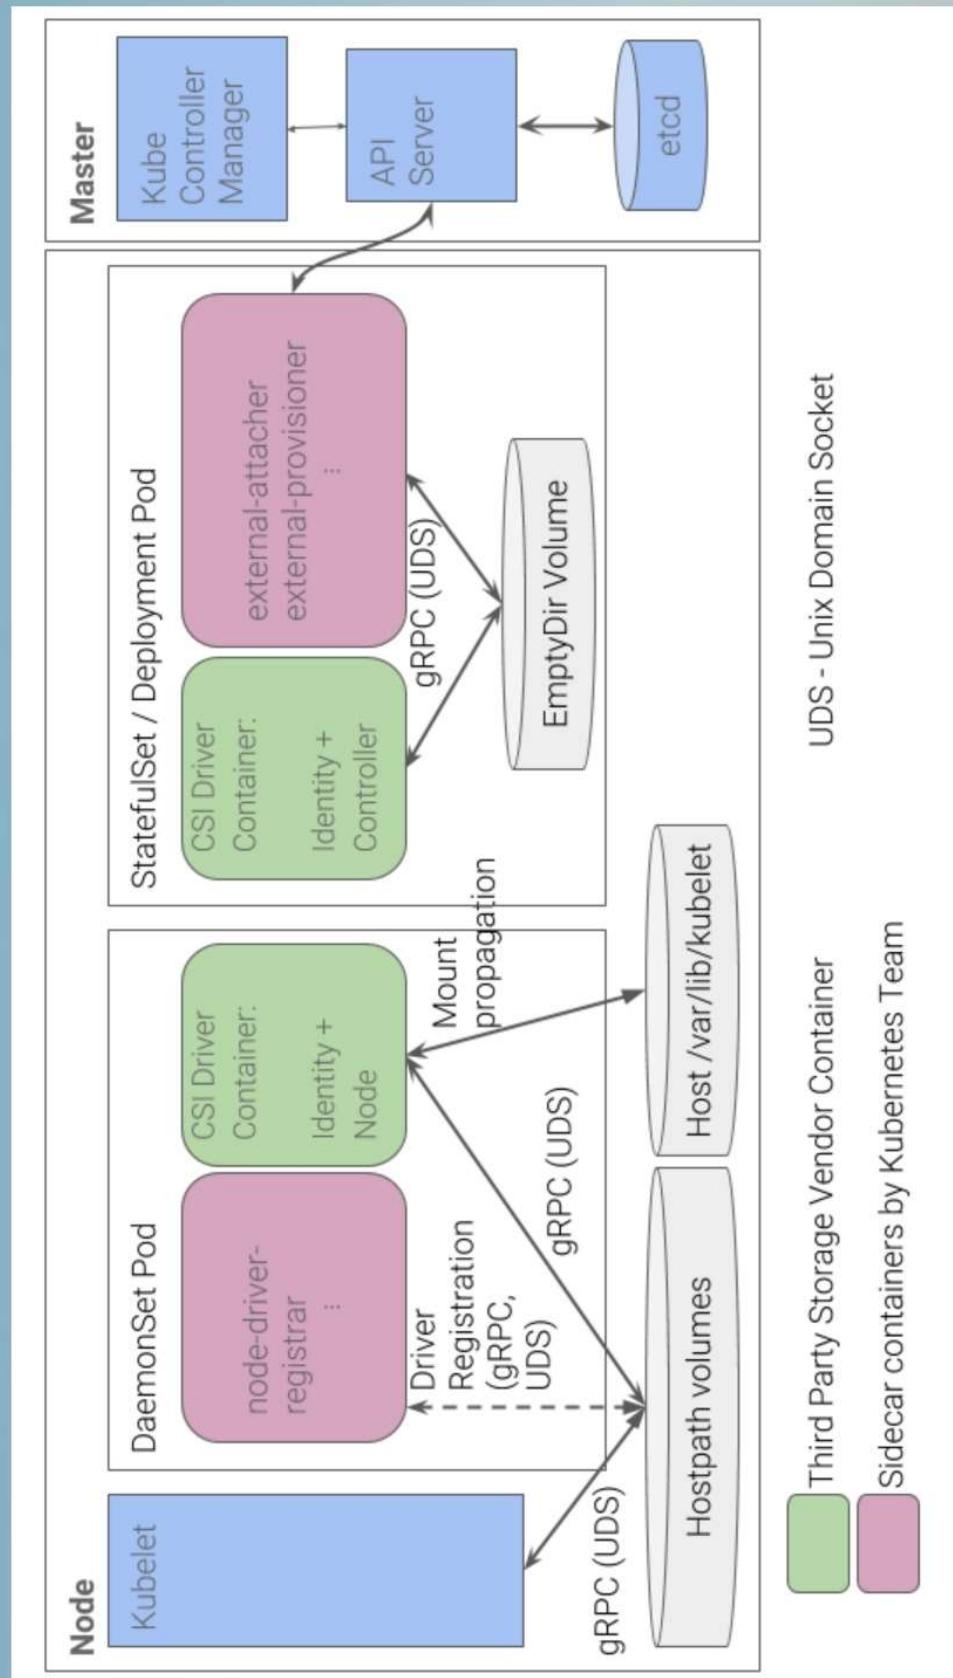
```
apiVersion: v1
kind: Pod
metadata:
  name: www
  labels:
    name: www
spec:
  containers:
    - name: www
      image: nginx:alpine
      ports:
        - containerPort: 80
          name: www
  volumeMounts:
    - name: www-persistent-storage
      mountPath: /usr/share/nginx/html
volumes:
  - name: www-persistent-storage
    persistentVolumeClaim:
      claimName: claim-http
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-0001
  labels:
    storage-context: customers
spec:
  capacity:
    storage: 2Gi
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
  selector:
    matchLabels:
      storage-context: customers
  nfs:
    server: 172.17.0.48
    path: /exports/data-0001
```

Recommended Mechanism for Deploying CSI Drivers on Kubernetes



CONTAINER
STORAGE
INTERFACE



<https://kubernetes-csi.github.io/docs/deploying.html>

<https://bit.ly/31T3lw>

Volumes

Containers, Pods, Nodes,
and Cluster come and go.
Volumes remain



SCENARIO

Kubernetes Fundamentals: Volumes and Mounts

An outward look at the file systems and volumes
available to your app.

Lab

Explore

- Setup a volume
- Declare volume access
- Connect Pod to volume



Probes

Checking Container and Pod Health

Lab

- **Liveness**
Are you even responding?
If not, Pod is killed (moo)



Probes

Checking Container and Pod Health

Lab

- **Liveness**
Are you even responding?
If not, Pod is killed (moo)

- **Readiness**

Are you available for traffic?
Service load balancer skips Pod until healthy



Probes

Checking Container and Pod Health

Lab

- **Liveness**

Are you even responding?
If not, Pod is killed (moo)

- **Readiness**

Are you available for traffic?
Service load balancer skips Pod until healthy

- **Startup (new)**

Geez, you're taking a long time to initialize!
Databases, complex models, chunky legacy apps



Probes

Checking Container and Pod Health



SCENARIO

Kubernetes Fundamentals: Health Checks with Probes

Explore the importance and difference between the three probe types for Pods

Lab

- **Liveness**

Are you even responding?
If not, Pod is killed (moo)

- **Readiness**

Are you available for traffic?
Service load balancer skips Pod until healthy

- **Startup (new)**

Geez, you're taking a long time to initialize!
Databases, complex models, clunky legacy apps



Probes

Checking Container and Pod Health



```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
    name: liveness-http
spec:
  containers:
    - name: liveness
      image: k8s.gcr.io/liveness
      args:
        - /server
      livenessProbe:
        httpGet:
          path: /healthz
          port: 8080
        failureThreshold: 3
        initialDelaySeconds: 45
        periodSeconds: 30
```

- HTTP Status codes
- TCP Connection
- gRPC, HTTP
- Exec, zero=pass, non zero=fail

Probes

Checking Container and Pod Health

Adjust these behaviors to match
expected container personality:

- Pass/fail logic
- Communication type
- initialDelaySeconds
- periodSeconds
- timeoutSeconds
- successThreshold
- failureThreshold



```
apiVersion: v1
kind: Pod
metadata:
  name: goproxy
  labels:
    app: goproxy
spec:
  containers:
    - name: goproxy
      image: k8s.gcr.io/goproxy:0.1
      ports:
        - containerPort: 8080
          readinessProbe:
            tcpSocket:
              port: 8080
            initialDelaySeconds: 5
            periodSeconds: 10
            livenessProbe:
              tcpSocket:
                port: 8080
            initialDelaySeconds: 15
            periodSeconds: 20
```

A Top Reason Why Clusters Fail

Your applications in containers consume CPU and memory.
Don't keep your resource needs a secret from Kubernetes.

Office workers on
floor 7 1/2

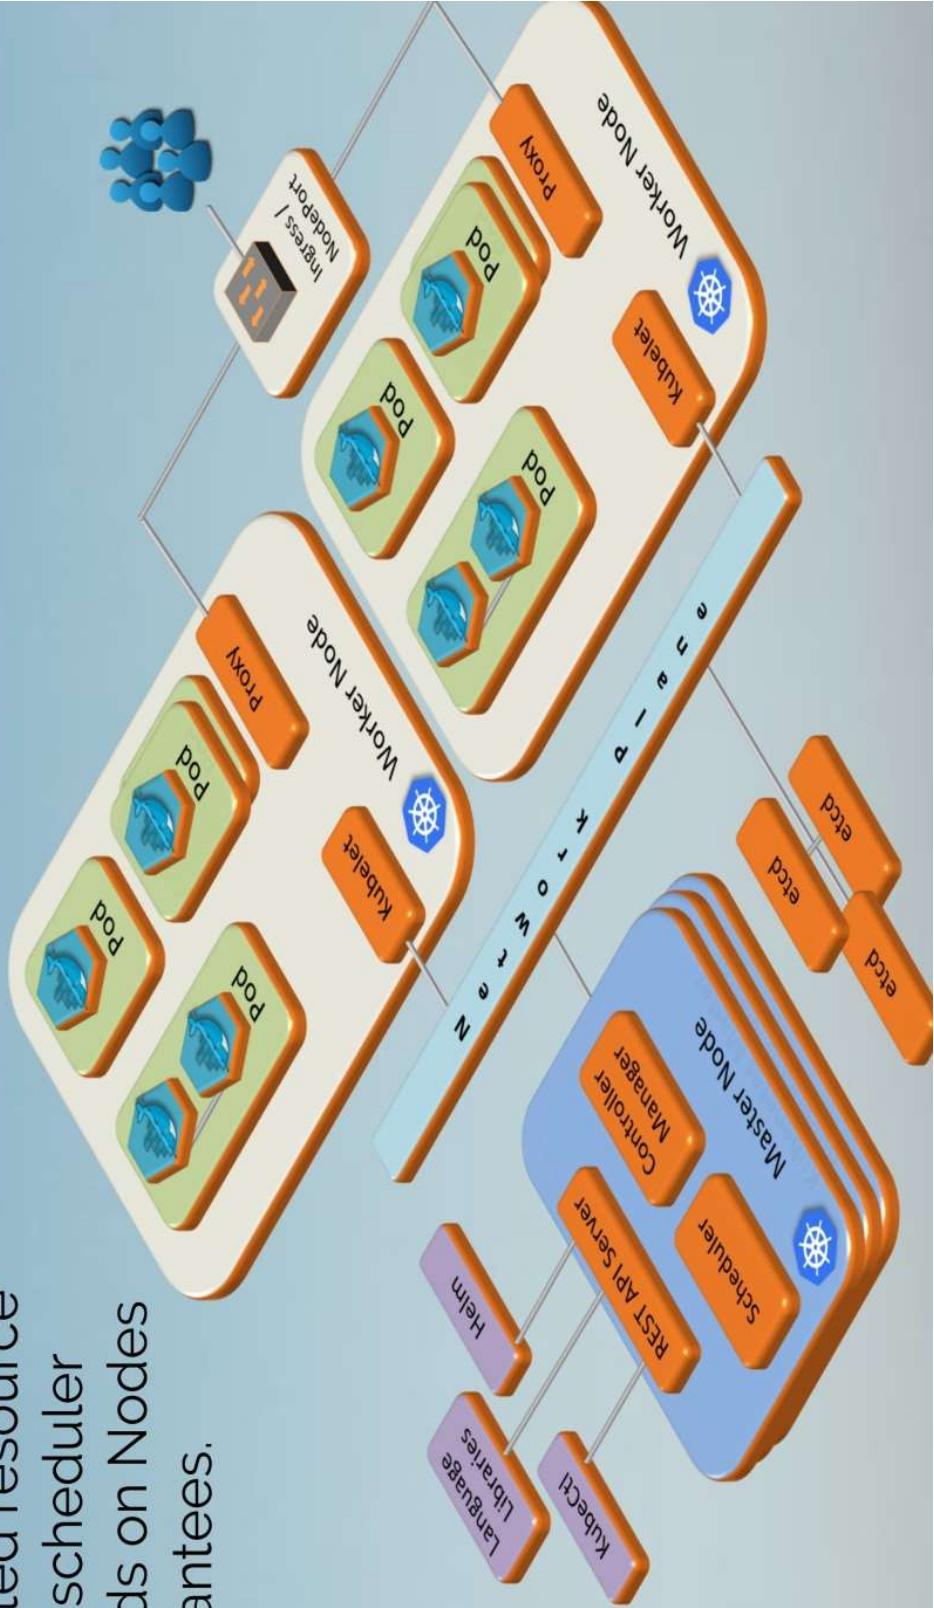
Being John Malkovich

1999 American fantasy
comedy film directed by
Spike Jonze and written
by Charlie Kaufman



CPU and Memory are limited resources

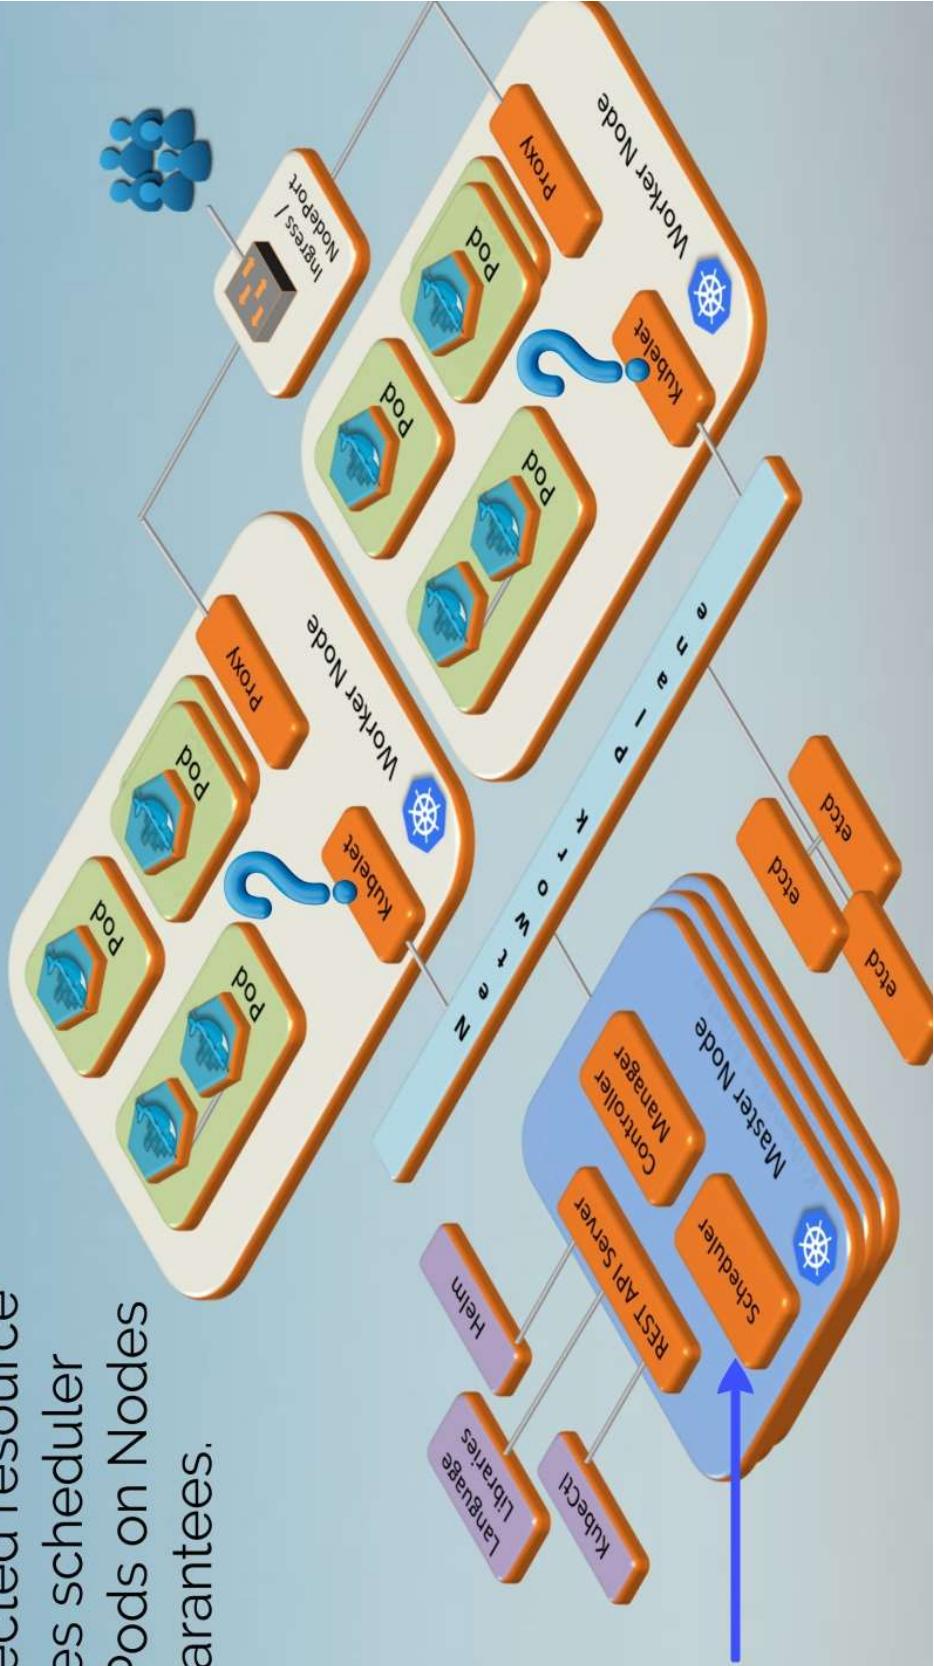
By providing expected resource needs, Kubernetes scheduler "bin packs" your Pods on Nodes with resource guarantees.



CPU and Memory are limited resources

By providing expected resource needs, Kubernetes scheduler "bin packs" your Pods on Nodes with resource guarantees.

Is there room?



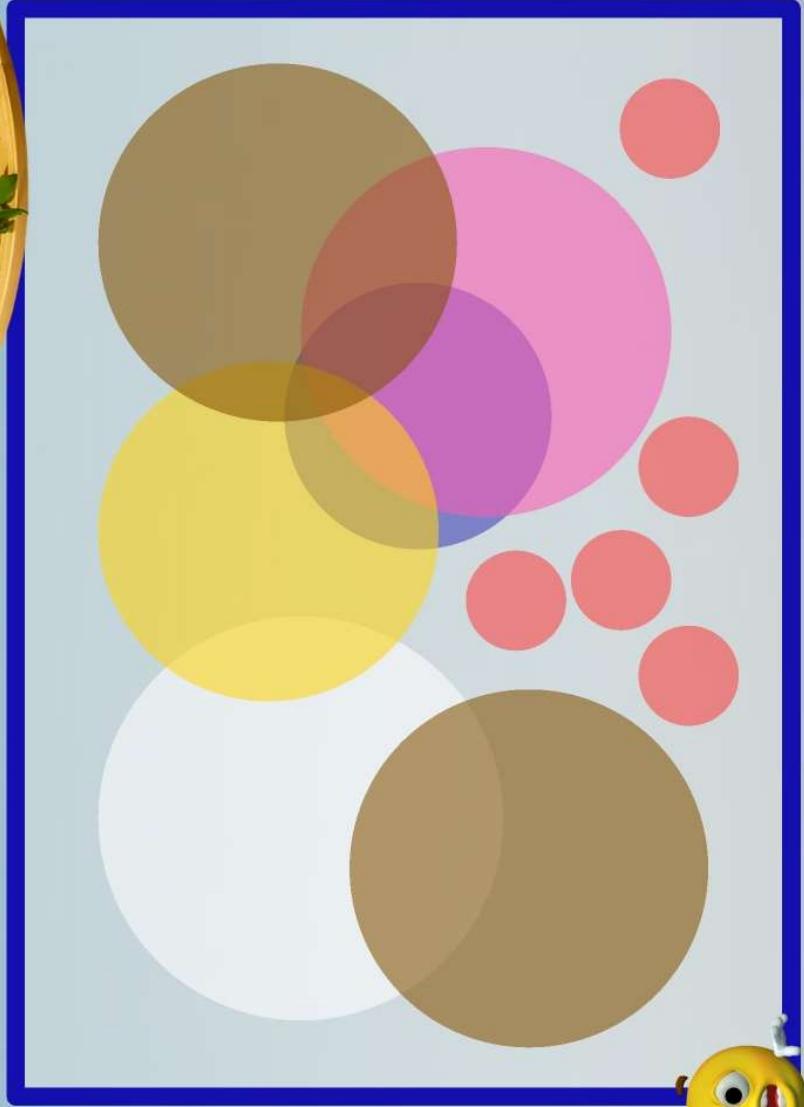
Quality of Service

BestEffort

- BestEffort
- Burstable
- Guaranteed

```
apiVersion: v1
kind: Pod
metadata:
  name: pack-em-in
spec:
  containers:
    - name: best-effort-example
      image: potential-beast:1.2.3
```

Dangerous - no CPU or memory requests
Rogue pods disrupt other pods and nodes



Quality of Service

Burstable

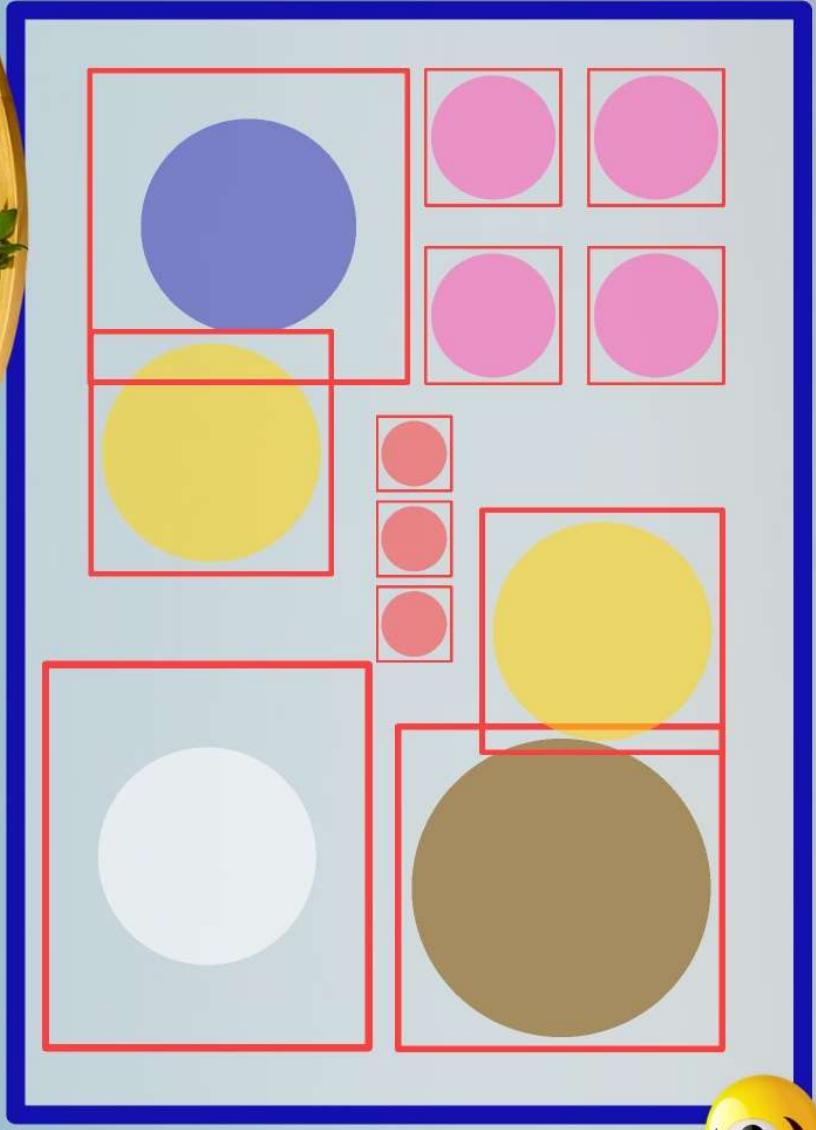
- BestEffort
- Burstable
- Guaranteed

```
apiVersion: v1
kind: Pod
metadata:
  name: pack-em-in
spec:
  containers:
    - name: best-effort-example
      image: potential-beast:1.2..3
  resources:
    limits:
      memory: "512Mi"
      cpu: "1024m"
    requests:
      memory: "256Mi"
      cpu: "600m"
```

Noisy Neighbor problem remains when too many bursts



Better, works with autoscalar



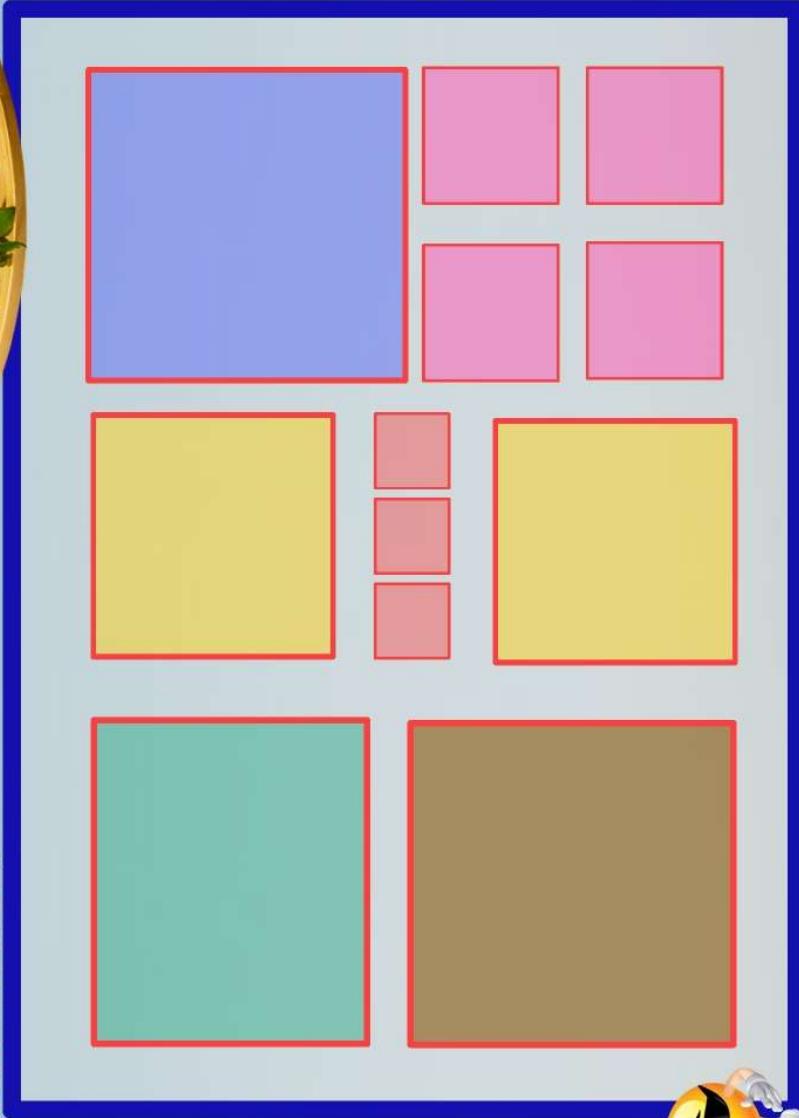
Quality of Service Guaranteed

- BestEffort
- Burstable
- Guaranteed



```
apiVersion: v1
kind: Pod
metadata:
  name: pack-em-in
spec:
  containers:
    - name: best-effort-example
      image: potential-beast:1.2.3
      resources:
        limits:
          memory: "256Mi"
          cpu: "1024m"
        requests:
          memory: "256Mi"
          cpu: "1024m"
```

Safest
Potentially over allocates and wastes (\$)



ResourceQuotas

- Apply ResourceQuota to limit available resource for:
 - Pods in Namespaces
 - Pods in an affinity group

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: pods-high
spec:
  hard:
    cpu: "1000"
    memory: 200Gi
    pods: "10"
  scopeSelector:
    matchExpressions:
      - operator: In
        scopeName: PriorityClass
        values: ["high"]
```



Hey, I Like You ... NOT!

3 native match making tricks through labels



Hey, I Like You ... NOT!

3 native match making tricks through labels

Node Affinity/Anti-affinity

Pods are marked to tell scheduler what **Nodes** they prefer, or not.



Hey, I Like You ... NOT!

3 native match making tricks through labels

Node Affinity/Anti-affinity

Pods are marked to tell scheduler what **Nodes** they prefer, or not.



Taints and Tolerations

Nodes are marked to tell scheduler what **Pods** they prefer, or not.

Hey, I Like You ... NOT!

3 native match making tricks through labels

Node Affinity/Anti-affinity

Pods are marked to tell scheduler what **Nodes** they prefer, or not.



Taints and Tolerations

Nodes are marked to tell scheduler what **Pods** they prefer, or not.

Pod Affinity/Anti-affinity

Pods are marked to tell scheduler what other **Pod** types they wish to be near (cloud zones), or not.

Hey, I Like You ... NOT!

3 native match making tricks through labels



Node Affinity/Anti-affinity

Pods are marked to tell scheduler what **Nodes** they prefer, or not.

```
kubectl label nodes <node-name>  
accelerator=nvidia-tesla-p100
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: cuda-vector-add  
spec:  
  restartPolicy: OnFailure  
  containers:  
    - name: cuda-vector-add  
      image: "k8s.gcr.io/cuda-vector-add:v0.1"  
  resources:  
    limits:  
      nvidia.com/gpu: 1  
  nodeSelector:  
    accelerator: nvidia-tesla-p100
```

Taints and Tolerations

Nodes are marked to tell scheduler what **Pods** they prefer, or not.

Pod Affinity/Anti-affinity

Pods are marked to tell scheduler what other **Pod** types they wish to be near (cloud zones), or not.